

Efficient and Scalable Bipartite Matching with Fast Beta Linkage (fabl)

Brian Kunding^{*}, Jerome Reiter^{*} and Rebecca C. Steorts[†]

Abstract. Recently, researchers have developed Bayesian versions of the Fellegi Sunter model for record linkage. These have the crucial advantage of quantifying uncertainty from imperfect linkages. However, current implementations of Bayesian Fellegi Sunter models are computationally intensive, making them challenging to use on larger-scale record linkage tasks. We propose a variation on Bayesian Fellegi Sunter models that we call fast beta linkage, or **fabl**. Specifically, we propose independent prior distributions over the matching space, allowing us to use hashing techniques that reduce computational overhead. This also allows us to complete pairwise record comparisons over large datasets through parallel computing and to reduce memory costs through a new technique called storage efficient indexing. Through simulations and two case studies, we show that **fabl** has markedly increased speed with minimal loss of accuracy.

BEKA/JERRY: Reviewer 2 wanted a specific reference to Sadinle in the abstract. I remember we actively decided not to do that. What do you think?

Keywords: entity resolution, data fusion, hashing, record linkage, Bayesian, distributed computing.

1 Introduction

In many data analysis tasks, analysts seek to identify duplicate records across two databases. This is an increasingly important task in “data cleaning,” and is used for inferential and predictive analyses in fields such as statistics, computer science, machine learning, political science, economics, precision medicine, official statistics, and others (Christen, 2012; Gutman et al., 2013; Dalzell and Reiter, 2018; Tang et al., 2020). In this article, we consider bipartite record linkage, which merges two databases that contain duplications across but not within the respective databases.

Many statistical record linkage methods are extensions of the seminal work of Fellegi and Sunter (1969) and Newcombe et al. (1959). Specifically, Fellegi and Sunter (1969) created comparison vectors for each pair of records in the data and independently classified each pair as a match or a non-match using a likelihood ratio test. Recent work in the statistical literature has extended this approach for a wide variety of applications (William E Winkler and Thibaudeau, 1990; Fair, 2004; Wagner et al., 2014; Gill and Goldacre, 2003; Enamorado et al., 2019; Aleshin-Guendel and Sadinle, 2022). Additionally, some methods model records directly (Steorts et al., 2016; Marchant et al., 2019; Betancourt et al., 2021), but in this paper, we build on the contributions to the comparison vector approach.

The independent pairwise matching assumption from Fellegi and Sunter (1969) is

^{*}Department of Statistical Science, Duke University, P.O. Box 90251, Durham, NC 27708, USA
© 0000 International Society for Bayesian Analysis DOI: 0000

[†]Departments of Statistical Science and Computer Science, Duke University, P.O. Box 90251, Durham, NC 27708, USA beka@stat.duke.edu

2 *Efficient and Scalable Bipartite Matching with Fast Beta Linkage (fabl)*

popular mainly for its mathematical simplicity, but can be unreasonable in practice. In many situations, we know that there are no duplications within a database, meaning that each record in one file should be linked with at most one other record in the other file. Thus, when the procedure results in many-to-one matches, some of these links must be false. Many extensions to Fellegi and Sunter (1969) resolve these false matches as a post-processing step (Jaro, 1989), but this model misspecification can still lead to poor results (Sadinle, 2017).

Alternatively, one can embed one-to-one matching requirements into the model specification itself (Gutman et al., 2013; Tancredi and Liseo, 2011), at an additional computational cost. Larsen (2005) employed a Metropolis-Hastings algorithm to only allow sampling matches that respected one-to-one assumptions, but such algorithms exhibit slow mixing due to the combinatorial nature of the constrained matching space. Fortunato (2010) used simulated annealing to target the space of matches permitted under the one-to-one constraint, but the method is computationally intensive and, to our knowledge, has not been applied on databases with more than 100 records. Sadinle (2017) proposed the Beta Record Linkage model (BRL), using a prior distribution over the space of bipartite matchings to strictly enforce one-to-one requirements throughout a Gibbs sampler. Additionally, he introduced a class of loss functions that allows for a flexible estimation of the linkage structure, such that the modeler can weight the relative importance of false positives and false negatives, and identify records pairings to be decided through clerical review. BRL was shown to work on larger tasks than previous one-to-one methods, but in our experience, it becomes slow when applied to files with more than a few thousand records.

In this paper, we propose fast beta linkage (**fabl**), which extends the BRL model for increased efficiency and scalability. Following the suggestion in Wortman (2019), we relax the one-to-one matching requirement of BRL and propose independent priors over the matching space. This allows us to (1) employ hashing techniques that speed up calculations and reduce computational costs, (2) compute the pairwise record comparisons over large datasets via parallel computing, and (3) reduce memory costs through what we call storage efficient indexing. We demonstrate **fabl** provides accurate estimation of the linkage structure and greatly enhanced speed. Open source software to use **fabl** in R is available through Github.

In what follows, Section 2 reviews the work of Fellegi and Sunter (1969) and Sadinle (2017). Section 3, proposes the **fabl** model, provides the Gibbs sampler for posterior inference, and shows the loss function used to calculate the Bayes estimate for the bipartite matching. Section 4 introduces the hashing technique and storage efficient indexing used to increase the speed of calculations and the scale of linkage tasks amenable to **fabl**. Sections 5 and 6 demonstrate the speed and accuracy of **fabl** through simulation studies and case studies of homicides from the El Salvadoran Civil War and the National Long Term Care Study.

2 Review of Prior Work

Consider two databases \mathbf{X}_1 and \mathbf{X}_2 containing records $\{x_{1i}\}_{i=1}^{n_1}$ and $\{x_{2j}\}_{j=1}^{n_2}$ respectively. Without loss of generality, denote files such that $n_1 \geq n_2$. We follow the convention established by Sadinle (2017) and say “record $i \in \mathbf{X}_1$ ” rather than the more compact x_{1i} in order to avoid double subscripts. In the context of bipartite matching, we assume that there are duplications across, but not within, each database. Under this framework, the set of matches across databases can be represented in two equivalent ways. First, we may use a matrix $\Delta \in \{0, 1\}^{n_1 \times n_2}$, where

$$\Delta_{ij} = \begin{cases} 1, & \text{if records } i \in \mathbf{X}_1 \text{ and } j \in \mathbf{X}_2 \text{ refer to the same entity;} \\ 0, & \text{otherwise.} \end{cases} \quad (1)$$

Though intuitive, this sparse matrix representation can become cumbersome for large linkage tasks. More compactly, bipartite matching also can be viewed as a labeling $\mathbf{Z} = (Z_1, \dots, Z_{n_2})$ for the records in database \mathbf{X}_2 such that

$$Z_j = \begin{cases} i, & \text{if records } i \in \mathbf{X}_1 \text{ and } j \in \mathbf{X}_2 \text{ refer to the same entity;} \\ n_1 + j, & \text{if records } j \in \mathbf{X}_2 \text{ does not have a match in database } \mathbf{X}_1. \end{cases} \quad (2)$$

We can go back and forth between the two using $\Delta_{ij} = I(Z_j = i)$, where $I(\cdot) = 1$ when the expression inside the parentheses is true, and $I(\cdot) = 0$ otherwise.

Denote the set of matches by $\mathbf{M} = \{(i, j) : i \in \mathbf{X}_1, j \in \mathbf{X}_2, \Delta_{ij} = 1\}$, and the set of non-matches by $\mathbf{U} = \{(i, j) : i \in \mathbf{X}_1, j \in \mathbf{X}_2, \Delta_{ij} = 0\}$. The record linkage task can be viewed as identifying the sets of \mathbf{M} and \mathbf{U} . We refer to record pairs that are estimated as matches as “links”, and record pairs that are estimated as non-matches as “non-links”.

Intuitively, matching records (those that refer to the same entity) should be similar; records that are not non-matching should not be similar. Fellegi and Sunter (1969) proposed encoding this is using a comparison vector γ_{ij} computed for each record pair (i, j) in $\mathbf{X}_1 \times \mathbf{X}_2$. Denote the number of criteria for comparing records by F , such that $\gamma_{ij} = (\gamma_{ij}^1, \gamma_{ij}^2, \dots, \gamma_{ij}^F)$. In most cases, γ_{ij} consists of one comparison for each feature shared between the two datasets. We refer to set of all comparison vectors γ_{ij} as Γ .

The simplest way to compare two records is to check for agreement or disagreement, and this is commonly used for categorical variables. For more complex measurements, we can take into account partial agreement to more richly characterize the comparison; for numerical data, we can use absolute difference, and for text data, we can use string distance metrics such as Levenstein or Jaro-Winkler distance (Cohen et al., 2003). We then can set thresholds that allow us to represent comparisons through discrete levels of disagreement (Bilenko and Mooney, 2006; Elmagarmid et al., 2007). Let $\mathcal{S}_f(i, j)$ denote a general similarity measure for feature f of records i and j , where the range of \mathcal{S}_f can be divided into $L_f + 1$ intervals denoted by $I_{f0}, I_{f1}, \dots, I_{fL_f}$. Following convention, I_{f0} represents the highest level of agreement (inclusive of complete agreement) and I_{fL_f} represents the highest level of disagreement (including complete disagreement). Thus,

4 *Efficient and Scalable Bipartite Matching with Fast Beta Linkage (fabl)*

we can construct comparison vectors in the following way: $\gamma_{ij}^f = l$ if $\mathcal{S}_f(i, j) \in I_{fl}$. The choice of I_{fl} are application specific, which we discuss in our simulation and case studies.

(I don't love this right here, but I don't really know where else to put it. But if it won't cause a problem with the reviewers, I'll just leave it.) In practice, it is not feasible to make all-to-all record comparisons as the computational complexity is of order $O(n_1 \times n_2)$. The most common solution is to utilize blocking, which places similar records into partitions, or "blocks," to reduce this computational burden (Steorts et al., 2014; Murray, 2016). In deterministic blocking, the modeler chooses a field thought to be highly reliable, and only compares records that agree on that field. The record linkage method is then applied independently across all blocks, which can be done in parallel for additional speed gains. However, blocking on an unreliable field can lead to missed matches, making this form of blocking undesirable in some situations (Steorts et al., 2014).

In this paper, we use hashing and a new technique called storage efficient indexing to increase the scale of the linkage tasks we can undertake without blocking. This is useful when there is no reliable blocking field available, or one desires estimates of model parameters for the entire sample in question. In practice, `fabl` can be combined with blocking, but all derivations, simulations, and case studies are presented here without blocking.

2.1 Fellegi-Sunter Model

Under the Fellegi Sunter model, we assume first, that comparison vectors are independent given their matching status, and second, that the match status of the record pairs are independent. This allows us to model the comparison data using mixture models in the following way:

$$\begin{aligned}\Gamma_{ij} &= \gamma_{ij} \mid \Delta_{ij} = 1 \stackrel{iid}{\sim} \mathcal{M}(\mathbf{m}), \\ \Gamma_{ij} &= \gamma_{ij} \mid \Delta_{ij} = 0 \stackrel{iid}{\sim} \mathcal{U}(\mathbf{u}), \\ \Delta_{ij} &\stackrel{iid}{\sim} \text{Bernoulli}(p),\end{aligned}\tag{3}$$

where \mathcal{M} and \mathcal{U} are the distributions for matching and non-matching record pairs, \mathbf{m} and \mathbf{u} are their respective sets of parameters, and p is the marginal probability that a record pair is a match. These \mathbf{m} and \mathbf{u} parameters are then used to calculate the likelihood ratio

$$w_{ij} = \frac{P(\gamma_{ij} \mid \Delta_{ij} = 1)}{P(\gamma_{ij} \mid \Delta_{ij} = 0)}.\tag{4}$$

This ratio will be large when there is strong evidence of the pair being a match, and small otherwise.

The Fellegi Sunter method then uses thresholds T_m and T_u such that all record pairs with $w_{ij} > T_m$ are declared as links, and all those with $w_{ij} < T_u$ are declared as non-links (with those such that $T_u < w_{ij} < T_m$ left undetermined and subject to clerical

review). This is a problem in practice as transitive closures can be violated. Since each w_{ij} is calculated independently, there can be situations in which both $w_{ij} > T_m$ and $w_{i'j} > T_m$, so both (i, j) and (i', j) are declared as links, even though such a matching is not allowed by assumption.

To address this issue, Jaro (1989) proposed solving the following linear sum assignment problem:

$$\begin{aligned} \max_{\Delta} \sum_{i=1}^{n_1} \sum_{j=1}^{n_2} w_{ij} \Delta_{ij} \quad \text{subject to} \quad & \Delta_{ij} \in \{0, 1\}; \\ & \sum_{i=1}^{n_1} \Delta_{ij} \leq 1, j = 1, 2, \dots, n_2; \text{ and} \\ & \sum_{j=1}^{n_2} \Delta_{ij} \leq 1, i = 1, 2, \dots, n_1. \end{aligned} \quad (5)$$

The solution to this problem, for which several simple algorithms exist, is a bipartite matching that maximizes the sum of the Fellegi Sunter weights among matched pairs. Jaro provided no theoretical justification for using this approach; however Sadinle (2017) recently showed that under certain conditions, (5) is the maximum likelihood estimate for bipartite matching.

2.2 Beta Record Linkage Model

To estimate a one-to-one matching without using a post-processing step, Sadinle (2017) incorporated this constraint directly into the model. He proposed the beta distribution for bipartite matching, given by

$$P(\mathbf{Z} | \alpha_\pi, \beta_\pi) = \frac{(n_1 - n_{12}(\mathbf{Z}))!}{n_1!} \frac{B(n_{12}(\mathbf{Z}) + \alpha_\pi, n_2 - n_{12}(\mathbf{Z}) + \beta_\pi)}{B(\alpha_\pi, \beta_\pi)},$$

where $B(\cdot, \cdot)$ represents the Beta function, hyperparameters α_π and β_π encode prior beliefs about the proportion of records in \mathbf{X}_2 that have matches in \mathbf{X}_1 , and $n_{12}(\mathbf{Z}) = \sum_{j=1}^{n_2} I(Z_j \leq n_1 + j)$ is the number of records in \mathbf{X}_2 that have a match in \mathbf{X}_1 . This prior induces a Gibbs sampler that strictly enforces one-to-one matching, removing previously matched records from the set of candidate records when sampling Z_j . This creates a dependency that makes the sampler inherently serial, which can be slow when working on large linkage tasks.

3 Fast Beta Linkage

Recall the compact representation of the matching structure $\mathbf{Z} = (Z_1, Z_2, \dots, Z_{n_2})$ with $Z_j \in \{1, 2, \dots, n_1, n_1 + j\}$, where $Z_j < n_1$ indicates a match and $Z_j = n_1 + j$ indicates a non-match. In contrast to the prior over the vector \mathbf{Z} from Sadinle (2017), we follow Wortman (2019) and use independent priors for each component Z_j . However, unlike

6 Efficient and Scalable Bipartite Matching with Fast Beta Linkage (fabl)

Wortman (2019) who proposes a flat prior for Z_j , we use a proper and flexible prior. We denote the fast beta prior as follows:

$$Z_j|\pi = \begin{cases} \frac{1}{n_1}\pi, & z_j \leq n_1; \\ 1 - \pi, & z_j = n_1 + j; \end{cases}$$

$$\pi \sim \text{Beta}(\alpha_\pi, \beta_\pi),$$

α_π and β_π are fixed hyperparameters that encode prior beliefs about the proportion of records in \mathbf{X}_2 that have matches in \mathbf{X}_1 . Note the prior from Wortman (2019) can be viewed as a special case of the fabl prior, with π fixed at the mean of a Beta $\left(1, \frac{1}{n_1}\right)$ random variable.

Remark. Intuitively, this set of priors says that record $j \in \mathbf{X}_2$ has some match in \mathbf{X}_1 with probability π , and that each record $i \in \mathbf{X}_1$ is equally likely to be that match. One non-informative choice might be $\pi \sim \text{Beta}(1, 1)$, which corresponds to a prior belief that non-matches and matches are equally likely, and another might be $\pi \sim \text{Beta}\left(1, \frac{1}{n_1}\right)$, which corresponds to a uniform prior on the labeling of \mathbf{Z} .

Furthermore, note that it is common to encounter missing information in datasets used for record linkage. Whether information is missing in record $i \in \mathbf{X}_1$ or $j \in \mathbf{X}_2$, the comparison vector γ_{ij} will have missing values. We handle this issue by assuming that this missingness occurs at random (MAR, by Little and Rubin (2002)). With this assumption, we can marginalize over the missing data, and do all computation simply using the observed data. To notate this, we use the indicator $I_{obs}(\cdot)$, which takes on the value 1 when the argument inside the parentheses is observed, and 0 otherwise. (I don't love this right here, but I don't really know where else to put it.)

Note that linkage in this setting is conducted at the record level, rather than at the record pair level as in the Fellegi Sunter model. That is, π under fabl estimates the proportion of records in \mathbf{X}_2 that have matches, while λ in the Fellegi Sunter model estimates the proportion of record pairs that are matches. We find π to be more interpretable parameter than λ in the bipartite case. In this setting, there are at most n_2 matching pairs out of $n_1 n_2$ total pairs, meaning that λ is bounded above by $\frac{1}{n_1}$ and tends towards 0 as the size of the linkage task grows. Additionally, while the Fellegi Sunter model makes $n_1 \times n_2$ independent matching decisions and BRL makes n_2 dependent matching decisions, fabl strikes a middle ground between the two, making n_2 independent matching decisions. As shown in Sections 5 and 6, this allows fabl to achieve the benefits from BRL while making efficiency gains possible by exploiting independence.

Importantly, this independence relaxes the one-to-one requirement from BRL; the Gibbs sampler under fabl does ensure that each record in \mathbf{X}_2 can be matched to at most one record in \mathbf{X}_1 , but allows for the possibility that multiple records in \mathbf{X}_2 match to the same record in \mathbf{X}_1 . To resolve many-to-one matchings and obtain a one-to-one estimate, we use the post processing procedure explained in Section 3.2.

With the fabl prior defined and our missingness assumptions addressed, we present the full model. Let $\mathbf{m} = (\mathbf{m}_1, \dots, \mathbf{m}_F)$, where $\mathbf{m}_f = (m_{f1}, \dots, m_{fL_f})$ and $m_{fl} =$

$P(\gamma_{ij}^f = l | Z_j = i)$ for all fields f and agreement levels l . Let \mathbf{u} parameters be defined analogously, with $u_{fl} = P(\gamma_{ij}^f = l | Z_j \neq i)$. Then, using the condensed notation $\Phi = \{\mathbf{m}, \mathbf{u}\}$, our model is:

$$\begin{aligned} \mathcal{L}(\mathbf{Z}, \Phi | \Gamma) &= \prod_{j=1}^{n_1} \prod_{i=1}^{n_2} \prod_{f=1}^F \prod_{l=1}^{L_f} \left[m_{fl}^{I(Z_j=i)} u_{fl}^{I(Z_j \neq i)} \right]^{I(\gamma_{ij}^f=l) I_{obs}(\gamma_{ij}^f)}, \\ \mathbf{m}_f &\sim \text{Dirichlet}(\alpha_{f0}, \dots, \alpha_{fL_f}), \\ \mathbf{u}_f &\sim \text{Dirichlet}(\beta_{f0}, \dots, \beta_{fL_f}), \\ Z_j | \pi &= \begin{cases} \frac{1}{n_1} \pi, & z_j \leq n_1; \\ 1 - \pi, & z_j = n_1 + j; \end{cases} \\ \pi &\sim \text{Beta}(\alpha_\pi, \beta_\pi), \end{aligned}$$

181 I actually don't like the Φ notation. It seems unnecessary. The only reason it's here
182 is because Sadinle used it. What do you guys think?

183 3.1 Gibbs Sampler

184 We initialize Φ from random draws from the prior distributions of \mathbf{m} and \mathbf{u} , and initialize
185 \mathbf{Z} to reflect no matches across datasets, that is $\mathbf{Z} = (n_1 + 1, \dots, n_1 + n_2)$. We then
186 sample Φ from its full conditional:

$$\begin{aligned} \mathbf{m}_f | \Gamma, \mathbf{Z} &\sim \text{Dirichlet}(\alpha_{f1}(\mathbf{Z}), \dots, \alpha_{fL_f}(\mathbf{Z})), \\ \mathbf{u}_f | \Gamma, \mathbf{Z} &\sim \text{Dirichlet}(\beta_{f1}(\mathbf{Z}), \dots, \beta_{fL_f}(\mathbf{Z})), \end{aligned}$$

187 where $\alpha_{fl}(\mathbf{Z}) = \alpha_{fl} + \sum_{i,j} I_{obs}(\gamma_{ij}^f) I(\gamma_{ij}^f = l) I(z_j = i)$, and $\beta_{fl}(\mathbf{Z}) = \beta_{fl} + \sum_{i,j} I_{obs}(\gamma_{ij}^f) I(\gamma_{ij}^f =$
188 $l) I(z_j \neq i)$.

Next, we sample \mathbf{Z} componentwise from the full conditionals for each Z_j :

$$p\left(Z_j^{(s+1)} = i | \Gamma, \Phi, \mathbf{Z}^{(s)}\right) \propto \begin{cases} w_{z_j, j}, & z_j \leq n_1; \\ n_1 \frac{n_2 - n_{12}(\mathbf{Z}) + \beta_\pi}{n_{12}(\mathbf{Z}) + \alpha_\pi}, & z_j = n_1 + j, \end{cases},$$

where

$$w_{ij} = \prod_{f=1}^F \prod_{l=1}^{L_f} \left(\frac{m_{fl}}{u_{fl}} \right)^{I(\gamma_{ij}^f=l) I_{obs}(\gamma_{ij}^f)}.$$

189 Derivations for these full conditionals can be found in Appendix 7.2.

190 3.2 Bayes Estimate

191 We calculate a Bayes estimate $\hat{\mathbf{Z}}$ for the linkage parameter \mathbf{Z} by assigning different
192 positive losses to different types of errors, and minimizing posterior expected loss. We

8 Efficient and Scalable Bipartite Matching with Fast Beta Linkage (fabl)

193 adopt the loss function proposed in [Sadinle \(2017\)](#), in which $\hat{Z}_j \in \{1, \dots, n_1, n_1 + j, R\}$,
 194 with R representing the option to leave the matching undetermined by the model.
 195 Specifically, we have

$$L(\hat{Z}_j, Z_j) = \begin{cases} 0, & \text{if } Z_j = \hat{Z}_j; \\ \theta_R, & \text{if } \hat{Z}_j = R; \\ \theta_{10}, & \text{if } Z_j \leq 1, \hat{Z}_j = n_1 + j; \\ \theta_{01}, & \text{if } Z_j = n_1 + j, \hat{Z}_j \leq n_1; \\ \theta_{11}, & \text{if } Z_j \leq n_1, \hat{Z}_j \leq n_1, Z_j \neq \hat{Z}_j. \end{cases}$$

196 Here, θ_R is the loss from not making a decision on the linkage status, θ_{10} is the loss
 197 from a false non-match, θ_{01} is the loss from a false match, and θ_{11} is the loss from the
 198 special case of a false match in which the record has a true match other than the one
 199 estimated by the model.

In general, we set $(\theta_{10}, \theta_{01}, \theta_{11}, \theta_R) = (1, 1, 2, \infty)$ inducing the intuitive decision rule

$$\hat{Z}_j = \begin{cases} i, & \text{if } P(Z_j = i|\Gamma) > \frac{1}{2}; \\ 0, & \text{otherwise.} \end{cases}$$

200 For a more in-depth explanation of this function and the induced Bayes estimate, see
 201 [Sadinle \(2017\)](#).

202 Since the Gibbs sampler does not strictly enforce one-to-one matching, it is possible
 203 for this Bayes estimate to link multiple records in \mathbf{X}_2 to one record in \mathbf{X}_1 . To achieve a
 204 Bayes estimate that fulfills one-to-one matching requirement, we simply minimize the
 205 expected loss subject to the constraint that $\hat{Z}_j \neq \hat{Z}_{j'}$ for all $j \neq j'$. In the event that we
 206 have two records j and j' such that both $P(Z_j = i|\Gamma) > \frac{1}{2}$ and $P(Z_{j'} = i|\Gamma) > \frac{1}{2}$, we
 207 accept the match with the highest posterior probability, and declare the other to have
 208 no match. A similar approach appears in the most probable maximal matching sets used
 209 by [Steorts et al. \(2016\)](#) to match records to latent entities.

4 Efficient and Scalable Implementation

210
 211 To increase the speed of the Gibbs sampler and expand the scale of the linkage tasks
 212 we can undertake, we propose hashing methods similar to those used by [Enamorado
 213 et al. \(2019\)](#) in the creation of **fastlink**, a fast and scalable implementation of the
 214 Fellegi Sunter method. The key insight is to recognize that record pairs contribute to
 215 posterior calculations only through the agreement pattern of the γ_{ij} vector. To make
 216 this more precise, let h_1, \dots, h_P denote the unique agreement patterns, and collect these
 217 unique patterns in the set $\mathcal{P} = \{h_1, \dots, h_P\}$. Here, $P = |\mathcal{P}|$ is the total number of
 218 unique agreement patterns. P is bounded above by $\prod_{f=1}^F L_f$ which does not depend
 219 on n_1 or n_2 . In this context, the integers $\{1, \dots, P\}$ serve as hashed values that encode
 220 the same information as the original vectors themselves. Whenever possible, we conduct
 221 calculations over these P agreement patterns, instead of the typical $n_1 \times n_2$ record pairs.

222 Additionally, we store “one-hot encodings” of these patterns rather than the original
 223 Fellegi Sunter comparison vectors, to aid in vectorized computations. See Appendix 7.4
 224 for details.

225 4.1 Data Representation, Hashing, and Storage

226 First, we hash record pairs of the same agreement pattern to unique integer values.
 227 Enamorado et al. (2019) accomplished this efficiently through the hashing function

$$h^*(i, j) = \sum_{f=1}^F I(\gamma_{ij}^f > 0) 2^{\gamma_{ij}^f + I(k>1) \sum_{e=1}^{k-1} (L_e - 1)}.$$

228 This function maps each agreement pattern to a unique integer, allowing us to store
 229 a scalar quantity instead of an entire vector for each record pair. For computational
 230 ease, we then map these integers to sequential integers from $\{1, \dots, P\}$ corresponding
 231 to the enumerated patterns in \mathcal{P} . When the (i, j) pair exhibits the p^{th} pattern, we say
 232 $(i, j) \in h_p$.

With all record pairs converted to hashed values, we can create a more compact nested list

$$\mathcal{R} = \left\{ \{r_{j_p}\}_{p=1}^P \right\}_{j=1}^{n_2},$$

where $r_{j_p} = \{i \in \mathbf{X}_1 | (i, j) \in h_p\}$ is a list of records in \mathbf{X}_1 that share agreement pattern p with record $j \in \mathbf{X}_2$. This representation is useful because it allows us to easily compute sufficient statistics

$$H_{j_p} = \sum_i I((i, j) \in h_p) = ||r_{j_p}||,$$

which is the number of records in \mathbf{X}_1 that share agreement pattern p with record $j \in \mathbf{X}_2$, and

$$H_p = \sum_{i,j} I((i, j) \in h_p) = \sum_j H_{j_p},$$

233 which is the number of comparison vectors in Γ that exhibit pattern p . For convenience,
 234 denote this set of sufficient statistics $\mathcal{H} = \{\{H_{j_p}\}, \{H_p\}\}$. As we will show in Section
 235 4.3, all posterior calculations are conducted with these sufficient statistics. Note that \mathcal{P} ,
 236 \mathcal{R} , and \mathcal{H} fully characterize the comparison matrix Γ with no loss of information, so we
 237 proceed with $\tilde{\Gamma} = \{\mathcal{P}, \mathcal{R}, \mathcal{H}\}$ for posterior inference.

238 4.2 Scaling to Large Linkage Tasks

239 The hashing procedure described above considerably reduces the memory needed to
 240 store the comparison information. That is, instead of storing $n_1 \times n_2$ comparison vectors,
 241 we only store the P unique vectors, and then $n_1 \times n_2$ scalar quantities relating record
 242 pairs to those vectors. However, even storing these $n_1 \times n_2$ scalar labels can become
 243 burdensome with large databases. Worse, the overwhelming majority of these labels
 244 relate to record pairs that are clear non-matches.

10 Efficient and Scalable Bipartite Matching with Fast Beta Linkage (fabl)

To address this issue, we propose storage efficient indexing (SEI). In standard indexing, one decides a priori certain criteria that all linked pairs must satisfy, and labels any record pairs that do not meet that criteria as non-links. For example, one might only consider pairs with a certain similarity score on a field deemed to be important (like first name), or only pairs that exactly match on a specified number of fields. While generally chosen to be quite loose, establishing these criteria requires knowledge of the problem and invites room for human error.

With SEI, however, we can reduce the comparison space, and its associated storage costs, while avoiding these drawbacks. Observe that all records $i \in \mathbf{X}_1$ that share agreement pattern p with $j \in \mathbf{X}_2$ have the same Fellegi Sunter weight w_{ij} , and therefore the same probability when sampling Z_j . Thus, we know that records $i \in r_{j_p}$ such that H_{j_p} is large are very unlikely to be sampled consistently enough to be deemed a match through the Bayes estimate. We know this regardless of the form of the agreement pattern itself, or its associated probabilities. Therefore, rather than store all of these unlikely record labels, we choose to store only a small number S of them in a new nested listed \mathcal{R}^{SEI} . instead of storing $n_1 \times n_2$ record labels, SEI allows us to store at most $n_2 \times P \times S$ labels, regardless of how large n_1 might be. **In practice, we recommend $S = 10$, as this reduces the number of stored indices for highly unlikely record pairings, but does is not likely to eliminate any of the indices for plausible matches. Choosing S too low, like $S = 1$ or $S = 2$, can concentrate undue mass on unlikely matches and distort linkage results.** Posterior calculations still attribute the appropriate weight to all records through the summary statistics in \mathcal{H} , and thus we can proceed with posterior inference through the memory reduced $\tilde{\Gamma}^{\text{SEI}} = \{\mathcal{P}, \mathcal{R}^{\text{SEI}}, \mathcal{H}\}$. For simplicity, we suppress “SEI” in future notation, and use $\tilde{\Gamma}$ without loss of generality.

For large data, we can partition the two datasets \mathbf{X}_1 and \mathbf{X}_2 into smaller chunks $\{\mathbf{X}_{1m}\}$ and $\{\mathbf{X}_{2n}\}$ for more manageable computations. On a single machine, we can read-in data sequentially, conduct hashing, compress information through SEI, and delete the original data from memory before continuing with the next chunk of data. With multiple cores or multiple machines, this can be done in parallel. Summary statistics from each pairwise chunk comparison can then be easily synthesized to recover sufficient statistics for the larger linkage task. Thus, the combination of hashing, SEI, and partitioning allows us to conduct linkage tasks over much larger datasets.

4.3 Efficient Posterior Inference

Calculated at the level of the record pairs, updating $\alpha_{fl}(\mathbf{Z})$ and $\beta_{fl}(\mathbf{Z})$ for each field and level in the linkage task constitutes $2 \times \sum_f L_f$ many summations over $n_1 \times n_2$ quantities. These are simple calculations, but become computationally burdensome when working on large linkage tasks.

Instead, we use one-hot encodings of the agreement patterns \mathcal{P} for more efficient calculations. Denote $H_p^m = \sum_{j=1}^{n_2} I((Z_j, j) \in h_p)$ to be the number of matching record pairs with agreement pattern p . It follows that the number of non-matching record pairs with agreement pattern p is $H_p^u = H_p - H_p^m$. Then, if α_0 and β_0 are vectors of length $\sum_f L_f$ of prior parameters for the \mathbf{m} and \mathbf{u} distributions respectively, the posterior

update becomes

$$\begin{aligned}\alpha(\mathbf{Z}) &= \alpha_0 + \sum_{p=1}^P H_p^m \times h_p, \\ \beta(\mathbf{Z}) &= \beta_0 + \sum_{p=1}^P H_p^u \times h_p.\end{aligned}$$

Although sampling Z_j from the full conditional provided in Section 3.1 is conceptually straightforward, it becomes computationally burdensome when n_1 is large. This is because sampling a value from n_1 many options with unequal weights requires normalizing the weights to probabilities, which has a computational cost that scales with n_1 . To speed up computation, we break this sampling step into two. First, we calculate the Fellegi Sunter weight w_p associated with each unique pattern and sample the agreement pattern between j and its potential link. Second, we sample the record label uniformly among records associated with that agreement pattern for that particular $j \in \mathbf{X}_2$. More concretely, define $h(Z_j)$ to be the agreement pattern between j and its potential match, and say $h(Z_j) = h_{P+1}$ when $Z_j = n_1 + j$. Then,

$$P\left(h\left(Z_j^{(s+1)}\right) = p \mid \tilde{\Gamma}\Phi, \mathbf{Z}^{(s)}\right) \propto \begin{cases} w_p \times H_{j_p}, & p \leq P; \\ n_1 \frac{n_2 - n_{12}(\mathbf{Z}) + \beta_\lambda}{n_{12}(\mathbf{Z}) + \alpha_\lambda}, & p = P + 1; \end{cases} \quad (6)$$

Since all posterior updates are governed by the agreement patterns of the record pairs rather than the record labels themselves, we complete the entire Gibbs sampler first at the level of the P agreement patterns with (6) above. After all iterations are complete, we can back-fill the records uniformly among records of the sampled agreement pattern through (6) below:

$$P\left(Z_j^{(s+1)} = i \mid h\left(Z_j^{(s+1)}\right) = p\right) = \begin{cases} \frac{1}{H_{j_p}}, & (i, j) \in h_p; \\ 0, & \text{otherwise.} \end{cases} \quad (7)$$

282 We emphasize the computational gains of this split sampler: the first step is a sample
283 from P options, where P does not scale with the size of the linkage task; and the second
284 step is sampling uniformly at random, which is computationally simple even for large
285 sets of candidate records. These changes drastically improve the speed of the sampler,
286 and each can be parallelized if desired for additional computational gains. We provide a
287 summary of the **fabl** method through pseudocode in Appendix 7.3. The computational
288 complexity of **fabl** is given in Lemma 1.

289 *Remark.* We note here that this split sampler can be adapted for use with BRL as well,
290 but would require updating $\{H_j\}$ to account for the previously matched records that no
291 are longer candidate matches for the current record. BRL would see considerable speed
292 up through the procedure, but computation time would still scale with the size of the
293 larger dataset.

294 **Lemma 1.** Recall that n_1 and n_2 are the number of records in \mathbf{X}_1 and \mathbf{X}_2 respectively, F
295 is the number of fields used for comparisons across records, and P is the number of patterns

12 Efficient and Scalable Bipartite Matching with Fast Beta Linkage (*fabl*)

that comparison vectors can exhibit. We assume B processors available for parallelization. Then, the overall computational complexity of *fabl* is $O(\frac{F}{B}n_1n_2) + O(n_2P)$.

Proof. To prove the computational complexity, we consider two steps — constructing the comparison vectors and the Gibbs sampler. The computational complexity of all pairwise comparisons across the two databases \mathbf{X}_1 and \mathbf{X}_2 is $O(Fn_1n_2)$. The hashing procedure for all pairwise comparisons is also of complexity $O(Fn_1n_2)$. With B processors available, we can split these computations across B equally sized partitions and compute these comparisons in parallel, so the complexity becomes $O(\frac{F}{B}n_1n_2)$. There are then trivial computational costs associated with synthesizing summary statistics across these partitions.

Without hashing, the computational complexity of updating the \mathbf{m} and \mathbf{u} parameters is $O(Fn_1n_2)$. However, by doing calculations over the agreement patterns rather than the individual records, hashing reduces the overall complexity to $O(P)$. The complexity of updating Z sequentially at the record level is $O(n_1n_2)$. With hashing, we split this sampling into two steps. First we sample the agreement pattern of the match with complexity $O(n_2P)$, and then we sample the record exhibiting that pattern with complexity $O(n_2)$. Thus the complexity of sampling Z is $O(n_2P)$. Since $P \ll n_1$ in all reasonable applications, we have reduced the complexity of the Gibbs sampler from $O(Fn_1n_2)$ under BRL to $O(n_2P)$ under *fabl*. In summary, the total computational complexity is $O(\frac{F}{B}n_1n_2) + O(n_2P)$. \square

5 Simulation Studies

We demonstrate the speed and accuracy of *fabl* as compared to BRL through several simulation studies.

5.1 Speed

In our first simulation, we generate comparison vectors from pre-specified distributions so that we can easily increase the size of the linkage problem. We use five binary comparisons with probabilities for matching and non-matching pairs shown in Table 1). \mathbf{u} parameters are chosen to emulate what is seen in practice for first name, last name, day of birth, and month of birth. The parameters for year would vary largely by context, so we choose them here to adjust the desired difficulty of the linkage task. We simulate these data for different values of n_1 and n_2 , always with $n_2/2$ records in common across datasets, and compare the run-time of *fabl* against BRL. Since we have 5 binary comparison fields, the number of unique patterns P is bounded above by $2^5 = 32$, a bound which is consistently attained in simulations with more records.

We note here that the implementation of BRL that we use is coded in C (Sadinle, 2017). In contrast, we use non-optimized code written in R for *fabl*. While this complicates comparisons, and indeed disfavors *fabl*, the computational speed gains for *fabl* are still evident, especially for larger sample sizes. Additionally, although *fabl* is amenable

	m		u	
	Agree	Disagree	Agree	Disagree
First Name	$\frac{19}{20}$	$\frac{1}{20}$	$\frac{1}{100}$	$\frac{99}{100}$
Last Name	$\frac{19}{20}$	$\frac{1}{20}$	$\frac{1}{100}$	$\frac{99}{100}$
Day	$\frac{20}{19}$	$\frac{1}{20}$	$\frac{1}{30}$	$\frac{29}{30}$
Month	$\frac{20}{19}$	$\frac{1}{20}$	$\frac{1}{12}$	$\frac{11}{12}$
Year	$\frac{20}{19}$	$\frac{1}{20}$	$\frac{1}{12}$	$\frac{11}{12}$


Table 1: Probabilities used for m and u distributions in simulation studies

Figure 1: Run-time for BRL and fabl to run 1000 Gibbs iterations, including hashing step for fabl, for increasing values of both n_1 and n_2 . We see near quadratic growth in run-time for BRL, and near linear growth for fabl.

334 to parallelization, this simulation was run on a single core. Implementing `fabl` in C++
 335 with parallelization for the hashing step and sampling the matching status of the record
 336 pairs should lead to even more computational gains.

337 In Figure 1, where we increase both n_1 and n_2 , BRL is faster than `fabl` for low sample
 338 sizes, but that `fabl` is significantly faster at handling larger data. In particular, run-time
 339 for BRL seems to grow quadratically (or linearly with the size of both \mathbf{X}_1 and \mathbf{X}_2) while
 340 run-time for `fabl` seems to grow linearly (in the size of only \mathbf{X}_2).

341 The above discussion suggests that for fixed n_2 , computation time for `fabl` should
 342 remain mostly constant with growing n_1 . In the simulation results displayed in Figure 2
 343 fixing $n_2 = 500$, we see linear growth for the run-time under BRL as n_1 increases, with
 344 much more static run-time under `fabl`. The slight increases in run-time for `fabl` are
 345 due primarily to the hashing step, which again can be run in parallel for large data. To
 346 illustrate that these trends are generalizeable to other specifications of the comparison
 347 vectors, we have included the run-time results for an additional simulation study, under
 348 different comparison vector settings, in Appendix 7.6.




Figure 2: Run-time for BRL and fabl to run 1000 Gibbs iterations, including hashing step for fabl, with increasing n_1 , and n_2 fixed at 500. We see linear growth in run-time for BRL, and near constant run-time for fabl.

5.2 Accuracy under Full Estimates

We replicate a simulation study from [Sadinle \(2017\)](#) to compare `fabl` against BRL on several synthetic datasets with varying amounts of error and overlap (the number of records in common across files). Following methods proposed by [Christen and Pudjijono \(2009\)](#) and [Christen and Vatsalan \(2013\)](#), clean records are first simulated from frequency tables for first name, last name, age, and occupation in Australia. Fields are then chosen for distortion uniformly at random. Names are subject to string insertions, deletions and substitutions, as well as common keyboard, phonetic, and optical recognition errors. Age and occupation are distorted through keyboard errors and missingness.

We create comparison vectors according to the default settings of the `compareRecords` function from the BRL package, shown in Table 2. Each simulation identifies matched individuals between two datasets, each with 500 records. We conduct linkage when matching records exhibit 1, 2, and 3 errors across the four fields, and when there are 50, 250, and 450 individuals in common across datasets. Under each of these settings, we have 100 pairs of simulated datasets in order to obtain uncertainty quantification on our performance metrics. We use flat priors for all m and u parameters, run the Gibbs Sampler for 1000 iterations, and discard the first 100 as burn-in. We calculate Bayes estimates of the linkage structure using the losses $\theta_R = \infty, \theta_{10} = 1, \theta_{01} = 1, \theta_{11} = 2$. Traceplots for parameters of interest for one example simulation are provided in Appendix

Fields	Similarity	Level of Disagreement			
		1	2	3	4
First and Last Name	Levenstein	0	(0, .25]	(.25, .5]	(.5, .1]
Age and Occupation	Binary	Agree	Disagree		

Table 2: Construction of comparison vectors for accuracy study with simulated datasets.




Figure 3: Posterior means and credible intervals for accuracy metrics under the replication of simulation study from [Sadinle \(2017\)](#). For each level of overlap and each level of error, we have 100 paired sets of 500 records. Thus this table represents results for 900 datasets. We see comparable performance for all levels of error and overlap.

368 7.5; they show no obvious concern over MCMC convergence.

369 We compare **fabl** to BRL in terms of recall, precision and F-measure, as defined in
 370 [Christen \(2012\)](#). Recall is the proportion of true matches found by the model, that is,
 371 $\sum_{j=1}^{n_2} I(\hat{Z}_j = Z_j, Z_j \leq n_1) / \sum_{j=1}^{n_2} I(Z_j \leq n_1)$. Precision is proportion the of links found
 372 by the model that are true matches, that is $\sum_{j=1}^{n_2} I(\hat{Z}_j = Z_j, Z_j \leq n_1) / \sum_{j=1}^{n_2} I(\hat{Z}_j \leq n_1)$.
 373 The F-measure balances the two metrics to provide an overall measure of accuracy, and
 374 is defined as $2 \times (\text{Recall} + \text{Precision}) / (\text{Recall} \times \text{Precision})$. In Figure 3, we see that
 375 the two methods have comparable performance at all levels of error and overlap. In
 376 the specific case of high error and low overlap, widely regarded as the most difficult
 377 linkage scenario, we see that **fabl** performs slightly worse on average; however, the
 378 overall accuracy level remains high.

5.3 Accuracy under Partial Estimates

By leaving $\theta_{10} = \theta_{01} = 1$ and $\theta_{11} = 2$, but setting $\theta_R = 0.1$, we allow the model to decline to decide a match for certain records, with nonassignment being 10% as costly as a false match. In this context, we are no longer focused on finding all true matches, but rather protecting against false matches. Thus, instead of recall, we use the negative predictive value (NPV), defined as the proportion of non-links that are actual non-matches. Mathematically, $\text{NPV} = \sum_{j=1}^{n_2} I(\hat{Z}_j = Z_j = n_1 + j) / \sum_{j=1}^{n_2} I(\hat{Z}_j = n_1 + j)$. We continue to use the precision, which is renamed the positive predictive value (PPV) in this context. Lastly, we also examine the rejection rate (RR), or how often the model declines to make a linkage decision, defined as $\text{RR} = \sum_{j=1}^{n_2} I(\hat{Z}_j = R) / n_2$. **To convey this information alongside NPV and PPV, for which values close to 1 indicate strong performance, we report the decision rate (DR), defined as $\text{DR} = 1 - \text{RR}$.**

In Figure 4, we see that **fabl** maintains equivalently strong PPV as **BRL** across all linkage settings. However, with high amounts of error, and thus fewer accurate and discerning fields of information, the rejection rate under **fabl** rises, leading to a decrease in NPV. Since **fabl** does not remove previously matched records from consideration for a new record, posterior probabilities of matches at times can be split across more records; in contrast, **BRL** is able to maintain higher confidence in matches in this setting. If one wishes to use partial estimates, **fabl** will possibly leave more linkages for the modeller to match by hand than would be left under **BRL**, but the decisions made by each method will have nearly equal accuracy.

6 Case Studies

In our first case study, we revisit data from the El Salvadoran Civil War analyzed by [Sadinle \(2017\)](#). Though the data files used in this case study are small, it shows how the computational complexity of **fabl** depends on the number of unique agreement patterns found in the data, and how significant computational gains can be achieved by simplifying the construction of the comparison vectors. In the second case study, we apply **fabl** to link records from the National Long Term Care Study (NLTCs), a larger linkage task that is not feasible in reasonable time under **BRL** with typical computing setups.

6.1 Civilian Casualties from the El Salvadoran Civil War

The country of El Salvador was immersed in civil war from 1980 to 1991. Throughout the time, several organizations attempted to document casualties of the conflict. When estimating the total number of casualties, one cannot simply sum the numbers recorded by each organization, as it is likely that the same individuals are recorded in multiple casualty lists. Instead, record linkage techniques must be used to merge datasets before analyzing the data ([Lum et al., 2013](#)). We utilize lists of casualties from the war, one collected by El Rescate - Tutela Regal (ERTL) and another from the Salvadoran Human Rights Commission (CDHES, by its acronym in Spanish).¹ The ERTL dataset consists

¹We thank the Human Rights Data Analysis Group (HRDAG) for granting access to this data.




Figure 4: Negative predictive value (NPV), positive predictive value (PPV), and decision rate (DR) on simulated datasets. We see poorer performance for **fabl** only in situations with high overlap.

418 of digitized denunciations that had been published throughout the conflict, and the
 419 CDHES dataset consists of casualties that had been reported directly to the organization
 420 (Howland, 2008; Ball, 2000). The ERTL required additional investigation before recording
 421 denunciations as human rights abuses, and reports to the CHDES were made shortly
 422 after the events occurred; thus, both datasets are thought to be fairly reliable.

423 There are several challenges with these data. First, both datasets have been automat-
 424 ically digitized, which inherently leads to some degree of typographical error. Second,
 425 the only fields recorded are given name, last name, date of death, and place of death. It
 426 is relatively common for a parent and child to share the same given name, resulting in
 427 indistinguishable records for two different individuals.

428 Following Sadinle (2017), we utilize records that have non-missing entries for given
 429 and last name, which results in $n_1 = 4420$ records in CHDES and $n_2 = 1323$ records in
 430 ERTL. We standardize names to account for common misspellings and use a modified
 431 Levenshtein distance when comparing names to account for the fact that second names
 432 are often omitted in Spanish. Place of death is recorded by municipality and department
 433 within that municipality; however, since department is missing in 95% of records in
 434 CHDES and 80% of records in ERTL, we exclude department from our analysis. Thus,
 435 we conduct record linkage using given name, last name, municipality, and day, month,
 436 and year of death. We again use flat priors for the \mathbf{m} and \mathbf{u} parameters.

Fields	Similarity	Level of Disagreement			
		1	2	3	4
First and Last Name	Modified Levenstein	0	(0, .25]	(.25, .5]	(.5, 1]
Year of Death	Absolute Difference	0	1	2	3+
Month of Death	Absolute Difference	0	1	2-3	4+
Day of Death	Absolute Difference	0	1-2	3-7	8+
Municipality	Binary	Agree	Disagree		

Table 3: Construction of comparison vectors for El Salvador data resembling original implementation from (Sadinle, 2018). This set up leads to 2048 possible agreement patterns in total.

Fields	Similarity	Level of Disagreement		
		1	2	3
First and Last Name	Modified Levenstein	0	(0, .25]	(.25, 1]
Year of Death	Binary	Agree	Disagree	
Month of Death	Binary	Agree	Disagree	
Day of Death	Absolute Difference	0	1	2+
Municipality	Binary	Agree	Disagree	

Table 4: Construction of comparison vectors for El Salvador for increased speed under **fabl**. This set up leads to 216 possible agreement patterns in total.

We initially followed the comparison vector constructions set by Sadinle (2017), using using four levels of agreement for each field, according to the thresholds provided in Table 3. This results in $4^5 \times 2 = 2048$ possible agreement patterns, with 1173 patterns realized in the data. However, after investigating the posterior distributions of the \mathbf{m} and \mathbf{u} parameters, we noticed that posterior distributions of several levels of the \mathbf{m} and \mathbf{u} parameters are nearly identical, suggesting that these levels are unnecessary.

Therefore we re-ran our analysis with fewer agreement levels for each field according to Table 4. Among the 216 possible agreement patterns, 159 are realized in the data. With this revised comparison specification, **fabl** ran in 61 seconds, approximately 4 times faster than the **BRL** run time of 239 seconds. The estimates of the \mathbf{m} parameters under each method are very similar, as shown in Figure 6. Traceplots for parameters of interest are provided in Appendix 7.7.

For completeness, we note that the more detailed comparison vector produced a run time of 240 seconds for **BRL**, and 261 seconds for **fabl**. Apparently, the number of patterns was sufficiently many that the computational savings from **fabl** does not overcome the inherent speed differences of C as opposed to R.

Through **fabl**, we arrive at a Bayes estimate of 179 individuals recorded in both datasets. We calculate posterior samples of the size of the overlap across files by finding the number of matches found in each iteration of the Gibbs sampler, and subtracting the number of matches that violate one-to-one matching. The posterior 95% credible interval for the overlap across files is (206, 238), indicating that the Bayes estimate identifies fewer matches than the Gibbs sampler identifies on average. This is because




Figure 5: Posterior distribution and Bayes estimate of overlap across the two files. We note they are quite similar under both methods.

a large number of records in ERTL have multiple plausible matches in CDHES; `fabl` recognizes that a match exists among the several options, but is unable to definitely declare a specific pair as a match in the Bayes estimate. We also compute a partial estimate of the linkage structure, using $\theta_{10} = \theta_{01} = 1$, $\theta_{11} = 2$, and $\theta_R = 0.1$ as in the simulation study in Section 5.2. Here, the Bayes estimate provides 136 matches of which the model is quite confident, and 175 records to verify manually. This means that after clerical review, the number of individuals replicated across datasets would fall in the interval (136, 311), encapsulating the posterior credible interval. More or fewer records could be identified for clerical review by decreasing or increasing θ_R .

We see similar results under `BRL`, with a Bayes estimate of 181 individuals recorded in both datasets, a posterior 95% credible interval of (211, 244), and a range of (140, 294) after the partial estimate and clerical review. See Figure 5 for a visual comparison of the Bayes estimate and posterior credible intervals for the two methods. We note that Bayes estimates falling outside of posterior credible intervals has occurred previously in the record linkage literature (Sadinle, 2017; Steorts et al., 2016), and remains a topic for future research.




Figure 6: Posterior estimates of m parameters with 95% credible intervals for El Salvador case study. They are quite similar across the two methods.

6.2 National Long Term Care Study

The National Long Term Care Study (NLTCS) is a longitudinal study tracking the health outcomes of Medicare recipients. The initial survey began in 1982, with follow-up surveys taken approximately every five years. For this study, we seek to link records over the $n_1 = 20485$ records from 1982 to the $n_2 = 17466$ records from 1989. In each iteration of the study, some participants are added to the lists, and others are lost due to loss of contact or death. The NLTCS data have longitudinal links, so that in reality one does not need to conduct record linkage. However, following the strategy in Guha et al. (2021) we break the longitudinal links and treat the data from 1982 and 1989 as stand-alone files.

We link records using sex, date of birth, and location using the thresholds shown in Table 5. Storing three comparison scores for each of $20485 \times 17466 \approx 400,000,000$ record pairs would require approximately 8GB of memory. Standard settings on a 16GB personal computer do not allow storage of an object of this size, and thus BRL is unable to perform this linkage task on such a machine. However, through the fabl framework, we compute comparisons over 30 smaller comparison tasks, hash results, and compress information through storage efficient indexing; the resulting data object is just 10 MB, approximately 0.1% of what is required for the raw comparisons. Constructing the comparisons sequentially took approximately 40 minutes, which could be reduced considerably through parallel computing.

We run a Gibbs sampler for 1000 iterations, taking about 235 seconds. As shown in 7, the Bayes estimate of the linkage structure consists of 9634 matches, with a 95% credible interval of (9581, 9740). Since we have access to the true linkage structure, we can calculate recall to be 0.89 and precision to be 0.98, resulting in an F-measure of 0.94. Traceplots indicate quick convergence, and are similar to those seen in Appendix 7.5 and 7.7

Fields	Similarity	Level of Disagreement		
		1	2	3
Sex	Binary	Agree	Disagree	
Year of Birth	Binary	Agree	Disagree	
Month of Birth	Binary	Agree	Disagree	
Day of Birth	Binary	Agree	Disagree	
Location	Custom	Same State and Office	Same State	Otherwise

Table 5: Construction of comparison vectors for NTLCS data

7 Conclusion

In this paper, we have proposed **fabl**, a Bayesian record linkage method that scales to large data sets. We have proven that the proposed hashing method and model assumptions allow for a linkage procedure whose computational complexity does not scale with the size of the larger dataset. This makes **fabl** computationally advantageous in many linkage scenarios, particularly when one datafile is substantially smaller than the other. We have demonstrated the speed and accuracy of **fabl** by replicating a simulation study and a case study in [Sadinle \(2017\)](#), and through an additional case study that is computationally infeasible under BRL. Open source software to use **fabl** in R is available through [Github](#).




Figure 7: Posterior distribution and Bayes estimate of overlap across years 1982 and 1989 of NLTC data.

References

- Aleshin-Guendel, S. and Sadinle, M. (2022), “Multifile Partitioning for Record Linkage and Duplicate Detection,” *JASA*. [1](#)
- Ball, P. (2000), “The Salvadoran Human Rights Commission: Data Processing, Data Representation, and Generating Analytical Reports,” in *Making the Case: Investigating Large Scale Human Rights Violations Using Information Systems and Data Analysis*, eds. P. Ball, H. F. Spirer, and L. Spirer, pp. 15–24, American Association for the Advancement of Science. [17](#)
- Betancourt, B., Sosa, J., and Rodríguez, A. (2021), “A Prior for Record Linkage Based on Allelic Partitions,” *forthcoming*. [1](#)
- Bilenko, M. and Mooney, R. (2006), “Riddle: Repository of Information on Duplicate Detection, Record Linkage, and Identity Uncertainty,” Online; retrieved July 29, 2020. [3](#)
- Christen, P. (2012), “A Survey of Indexing Techniques for Scalable Record Linkage and Deduplication,” *IEEE Transactions on Knowledge and Data Engineering*, 24, 1537–1555. [1](#), [15](#)
- Christen, P. and Pudjijono, A. (2009), *Accurate Synthetic Generation of Realistic Personal Information*, Springer Berlin Heidelberg, Berlin, Heidelberg. [14](#)
- Christen, P. and Vatsalan, D. (2013), “Flexible and Extensible Generation and Corruption of Personal Data,” in *Proceedings of the 22nd ACM International Conference on Information and Knowledge Management*, CIKM ’13, p. 1165–1168, New York, NY, USA, Association for Computing Machinery. [14](#)
- Cohen, W. W., Ravikumar, P., and Fienberg, S. E. (2003), “A Comparison of String Distance Metrics for Name-Matching Tasks,” in *Proceedings of the 2003 International Conference on Information Integration on the Web*, p. 73–78, AAAI Press. [3](#)
- Dalzell, N. M. and Reiter, J. P. (2018), “Regression Modeling and File Matching Using Possibly Erroneous Matching Variables,” *Journal of Computational and Graphical Statistics*, 0, 1–11. [1](#)
- Elmagarmid, A. K., Ipeirotis, P. G., and Verykios, V. S. (2007), “Duplicate Record Detection: A Survey,” *IEEE Transactions on Knowledge and Data Engineering*, 19, 1–16. [3](#)
- Enamorado, T., Fifield, B., and Imai, K. (2019), “Using a Probabilistic Model to Assist Merging of Large-Scale Administrative Records,” *American Political Science Review*, 113, 353–371. [1](#), [8](#), [9](#)
- Fair, M. (2004), “Generalized Record Linkage System—Statistics Canada’s Record Linkage Software,” *Austrian Journal of Statistics*, 33, 37–53. [1](#)
- Fellegi, I. P. and Sunter, A. B. (1969), “A Theory for Record Linkage,” *Journal of the American Statistical Association*, 64, 1183–1210. [1](#), [2](#), [3](#)

- Fortunato, S. (2010), “Community Detection in Graphs,” *Physics Reports*, 486, 75–174. [2](#)
- Gill, L. and Goldacre, M. (2003), “English National Record Linkage of Hospital Episode Statistics and Death Registration Records,” *Report to the Department of Health*. [1](#)
- Guha, S., Reiter, J. P., and Mercatanti, A. (2021), “Bayesian Causal Inference with Bipartite Record Linkage,” *Bayesian Analysis*, p. in press. [20](#)
- Gutman, R., Afendulis, C. C., and Zaslavsky, A. M. (2013), “A Bayesian Procedure for File Linking to Analyze End-of-Life Medical Costs,” *Journal of the American Statistical Association*, 108, 34–47. [1](#), [2](#)
- Howland, T. (2008), “How El Rescate, a Small Nongovernmental Organization, Contributed to the Transformation of the Human Rights Situation in El Salvador,” *Human Rights Quarterly*, 30, 703–757. [17](#)
- Jaro, M. A. (1989), “Advances in Record-Linkage Methodology as Applied to Matching the 1985 Census of Tampa, Florida,” *Journal of the American Statistical Association*, 84, 414–420. [2](#), [5](#)
- Larsen, M. D. (2005), “Advances in Record Linkage Theory: Hierarchical Bayesian Record Linkage Theory,” in *Proceedings of the Joint Statistical Meetings, Section on Survey Research Methods*, pp. 3277–3284, The American Statistical Association. [2](#)
- Little, R. and Rubin, D. (2002), *Statistical Analysis with Missing Data*, Wiley, Hoboken, New Jersey. [6](#)
- Lum, K., Price, M. E., and Banks, D. (2013), “Applications of Multiple Systems Estimation in Human Rights Research,” *The American Statistician*, 67, 191–200. [16](#)
- Marchant, N. G., Steorts, R. C., Kaplan, A., Rubinstein, B. I. P., and Elazar, D. N. (2019), “d-blink: Distributed End-to-End Bayesian Entity Resolution,” *arXiv e-prints*, arxiv:1909.06039. [1](#)
- Murray, J. S. (2016), “Probabilistic Record Linkage and Deduplication after Indexing, Blocking, and Filtering,” *Journal of Privacy and Confidentiality*, 7, 3–24. [4](#)
- Newcombe, H. B., Kennedy, J. M., Axford, S. J., and James, A. P. (1959), “Automatic Linkage of Vital Records,” *Science*, 130, 954–959. [1](#)
- Sadinle, M. (2017), “Bayesian Estimation of Bipartite Matchings for Record Linkage,” *Journal of the American Statistical Association*, 112, 600–612. [2](#), [3](#), [5](#), [8](#), [12](#), [14](#), [15](#), [16](#), [17](#), [18](#), [19](#), [21](#), [28](#)
- Sadinle, M. (2018), “Bayesian Propagation of Record Linkage Uncertainty Into Population Size Estimation of Human Rights Violations,” *The Annals of Applied Statistics*, 12, 1013–1038. [18](#)
- Steorts, R. C., Ventura, S. L., Sadinle, M., and Fienberg, S. E. (2014), “A Comparison of Blocking Methods for Record Linkage,” in *Privacy in Statistical Databases*, Lecture Notes in Computer Science, pp. 253–268, Springer, Cham. [4](#)

- 587 Steorts, R. C., Hall, R., and Fienberg, S. E. (2016), “A Bayesian Approach to Graphical
588 Record Linkage and Deduplication,” *Journal of the American Statistical Association*,
589 111, 1660–1672. [1](#), [8](#), [19](#)
- 590 Tancredi, A. and Liseo, B. (2011), “A Hierarchical Bayesian Approach to Record Linkage
591 and Size Population Problems,” *Annals of Applied Statistics*, 5, 1553–1585. [2](#)
- 592 Tang, J., Reiter, J., and Steorts, R. C. (2020), “Bayesian Modeling for Simultaneous
593 Regression and Record Linkage,” in *Privacy in Statistical Databases*, Lecture Notes in
594 Computer Science, Springer, Cham. [1](#)
- 595 Wagner, D., Lane, M., et al. (2014), “The Person Identification Validation System (PVS):
596 Applying the Center for Administrative Records Research and Applications’ (CARRA)
597 Record Linkage Software,” Tech. rep., Center for Economic Studies, US Census Bureau.
598 [1](#)
- 599 William E Winkler and Thibaudeau, Y. (1990), “An Application of the Fellegi-Sunter
600 Model of Record Linkage to the 1990 US Decennial Census,” *Research Report*, pp.
601 1–22. [1](#)
- 602 Wortman, J. P. H. (2019), “Record Linkage Methods with Applications to Causal
603 Inference and Election Voting Data,” Ph.D. thesis, Duke University. [2](#), [5](#), [6](#)

Appendix

7.1 Summary of Notation

Symbol	Description
$\mathbf{X}_1, \mathbf{X}_2$	datasets
$i \in 1, \dots, n_1$	index over records in \mathbf{X}_1
$j \in 1, \dots, n_2$	index over records in \mathbf{X}_2
$f \in 1, \dots, F$	index over fields used for comparisons
$l \in 1, \dots, L_f$	index over agreement levels for field f
n_{12}	number of entities in common between \mathbf{X}_1 and \mathbf{X}_2
γ_{ij}	comparison vector for records $i \in \mathbf{X}_1$ and $j \in \mathbf{X}_2$
$Z_j = i$	records $i \in \mathbf{X}_1$ and $j \in \mathbf{X}_2$ match
$Z_j = n_1 + j$	record $j \in \mathbf{X}_2$ has no match in \mathbf{X}_1
m_{fl}	$P(\gamma_{ij}^f = l Z_j = i)$
u_{fl}	$P(\gamma_{ij}^f = l Z_j \neq i)$
π	probability that a record in \mathbf{X}_2 has a match in \mathbf{X}_1
r_{jp}	list of records in \mathbf{X}_1 that share agreement pattern p with record $j \in \mathbf{X}_2$
h_p	one hot encoding of agreement pattern p
H_{jp}	number of records in \mathbf{X}_1 that share agreement pattern p with record $j \in \mathbf{X}_2$
H_p	number of comparison vectors that exhibit pattern p

7.2 Derivations of Full Conditionals

The \mathbf{m} and \mathbf{u} parameters are updated through standard multinomial-Dirichlet distributions. For a particular $m_{f\ell}$, we have

$$\begin{aligned}
 p(m_{fl} | \Gamma, \mathbf{Z}) &\propto \prod_{i=1}^{n_1} \prod_{j=1}^{n_2} m_{fl}^{I(Z_j=i)I(\gamma_{ij}^f=l)I_{obs}(\gamma_{ij}^f)} \times m_{fl}^{\alpha_{fl}-1} \\
 &= m_{fl}^{\alpha_{fl}(\mathbf{Z})-1}
 \end{aligned}$$

where $\alpha_{fl}(\mathbf{Z}) = \alpha_{fl} + \sum_{i,j} I_{obs}(\gamma_{ij}^f)I(\gamma_{ij}^f = l)I(z_j = i)$. Analogous procedures lead to the posterior distribution $p(u_{fl} | \Gamma, \mathbf{Z}) \propto u_{fl}^{\beta_{fl}(\mathbf{Z})-1}$, where $\beta_{fl}(\mathbf{Z}) = \beta_{fl} + \sum_{i,j} I_{obs}(\gamma_{ij}^f)I(\gamma_{ij}^f = l)I(z_j \neq i)$. Thus for the vectors of parameters \mathbf{m}_f and \mathbf{u}_f , we have

$$\begin{aligned}
 \mathbf{m}_f | \mathbf{Z}, \Gamma &\sim \text{Dirichlet}(\alpha_{f1}(\mathbf{Z}), \dots, \alpha_{fL_f}(\mathbf{Z})) \\
 \mathbf{u}_f | \mathbf{Z}, \Gamma &\sim \text{Dirichlet}(\beta_{f1}(\mathbf{Z}), \dots, \beta_{fL_f}(\mathbf{Z})).
 \end{aligned}$$

In order to define the full conditional for \mathbf{Z} , we must provide the posterior distribution for π , and provide a more tractable expression for the portion of the likelihood relevant to a particular record $j \in \mathbf{X}_2$. Since π encodes the rate of matching across the two datasets, the posterior distribution $p(\pi | \Gamma, \mathbf{Z}, \Phi, \alpha_\pi, \beta_\pi)$ depends only on the number of

links $n_{12}(\mathbf{Z}) = \sum_{i=1}^{n_2} I(z_j < n_1 + j)$ encoded by \mathbf{Z} (and hyperparameters). Thus we use $p(\pi|\mathbf{Z}, \alpha_\pi, \beta_\pi)$, and have

$$\begin{aligned} p(\pi|\mathbf{Z}, \alpha_\pi, \beta_\pi) &\propto p(\mathbf{Z}|\pi)p(\pi) \\ &\propto \pi^{n_{12}(\mathbf{Z})}(1-\pi)^{n_2-n_{12}(\mathbf{Z})}\pi^{\alpha_\pi-1}(1-\pi)^{\beta_\pi-1} \\ &\propto \pi^{n_{12}(\mathbf{Z})+\alpha_\pi-1}(1-\pi)^{n_1-n_{12}(\mathbf{Z})+\beta_\pi-1}. \end{aligned}$$

Thus $\pi^{(s+1)}|\mathbf{Z}^{(s+1), \alpha_\pi, \beta_\pi}$ has a $\text{Beta}(n_{12}(\mathbf{Z}) + \alpha_\pi, n_2 - n_{12}(\mathbf{Z}) + \beta_\pi)$ distribution.

Let $\Gamma_{\cdot j}$ denote the set of n_1 comparison vectors with $j \in \mathbf{X}_2$. We have

$$\begin{aligned} p(\Gamma_{\cdot j}|Z_j = z_j, \Phi) &\propto \prod_{i=1}^{n_1} \left[\prod_{f=1}^F \prod_{l=1}^{L_f} m_{fl}^{I(Z_j=i)} u_{fl}^{I(Z_j \neq i)} \right]^{I(\gamma_{ij}^f=l)I_{obs}(\gamma_{ij}^f)} \\ &\propto \prod_{i=1}^{n_1} \frac{\left[\prod_{f=1}^F \prod_{l=1}^{L_f} m_{fl}^{I(Z_j=i)} u_{fl}^{I(Z_j \neq i)} \right]^{I(\gamma_{ij}^f=l)I_{obs}(\gamma_{ij}^f)}}{\prod_{f=1}^F \prod_{l=1}^{L_f} u_{fl}^{I(\gamma_{ij}^f=l)I_{obs}(\gamma_{ij}^f)}} \\ &\propto \prod_{i=1}^{n_1} \left(\prod_{f=1}^F \prod_{l=1}^{L_f} \frac{m_{fl}}{u_{fl}} \right)^{I(z_j=i)I(\gamma_{ij}^f=l)I_{obs}(\gamma_{ij}^f)} \\ &= \begin{cases} w_{z_j, j} & z_j \leq n_1; \\ 1 & z_j = n_1 + j, \end{cases} \end{aligned}$$

where

$$w_{ij} = \prod_{f=1}^F \prod_{l=1}^{L_f} \left(\frac{m_{fl}}{u_{fl}} \right)^{I(\gamma_{ij}^f=l)I_{obs}(\gamma_{ij}^f)}.$$

With the likelihood in this form, we derive an expression for the posterior distribution of \mathbf{Z} . Because we sample each Z_j independently of all other $Z_{j'}$, we use only the full conditional for an individual Z_j :

$$\begin{aligned} p\left(Z_j^{(s+1)}|\Gamma_{\cdot j}, \Phi, \mathbf{Z}^{(s)}, \pi\right) &\propto p(\Gamma_{\cdot j}|Z_j^{(s+1)}, \Phi)P(Z_j^{(s+1)}|\pi) \\ &\propto \left(\sum_{i=1}^{n_1} w_{z_j, j} I(z_j = i) + I(z_j = n_1 + j) \right) \left(\pi \sum_{i=1}^{n_1} \frac{1}{n_1} I(z_j = i) + (1-\pi)I(z_j = n_1 + j) \right) \\ &= \frac{\pi}{n_1} \sum_{i=1}^{n_1} w_{z_j, j} I(z_j = i) + (1-\pi)I(z_j = n_1 + j). \\ &= \begin{cases} \frac{\pi}{n_1} w_{z_j, j} & z_j \leq n_1; \\ 1 - \pi & z_j = n_1 + j. \end{cases} \end{aligned}$$

28 *Efficient and Scalable Bipartite Matching with Fast Beta Linkage (fabl)*

611 For more direct comparability with the method from [Sadinle \(2017\)](#), we avoid
 612 sampling π directly, and instead integrate over it in the final full conditional for Z_j .

$$\begin{aligned}
 p\left(Z_j^{(s+1)} = i | \Gamma, \Phi, \mathbf{Z}^{(s)}\right) &= \int_{\pi} p\left(Z_j^{(s+1)} | \Gamma, \Phi, \mathbf{Z}^{(s)}, \pi\right) p\left(\pi | \mathbf{Z}^{(s)}\right) d\pi \\
 &= \int_{\pi} \left[\frac{\pi}{n_1} w_{z_j, j} I(z_j \leq n_1) + (1 - \pi) I(z_j = n_1 + j) \right] p\left(\pi | \mathbf{Z}^{(s)}\right) d\pi \\
 &= \frac{\int_{\pi} \pi p\left(\pi | \mathbf{Z}^{(s)}\right) d\pi}{n_1} w_{z_j, j} I(z_j \leq n_1) + \left(1 - \int_{\pi} \pi p\left(\pi | \mathbf{Z}^{(s)}\right) d\pi\right) I(z_j = n_1 + j) \\
 &= \frac{n_{12}(\mathbf{Z}) + \alpha_{\pi}}{n_1(n_2 + \alpha_{\pi} + \beta_{\pi})} w_{z_j, j} I(z_j \leq n_1) + \frac{n_2 - n_{12}(\mathbf{Z}) + \beta_{\pi}}{n_2 + \alpha_{\pi} + \beta_{\pi}} I(z_j = n_1 + j) \\
 &\propto w_{z_j, j} I(z_j \leq n_1) + n_1 \frac{n_2 - n_{12}(\mathbf{Z}) + \beta_{\pi}}{n_{12}(\mathbf{Z}) + \alpha_{\pi}} I(z_j = n_1 + j) \\
 &\propto \begin{cases} w_{z_j, j} & z_j \leq n_1; \\ n_1 \frac{n_2 - n_{12}(\mathbf{Z}) + \beta_{\pi}}{n_{12}(\mathbf{Z}) + \alpha_{\pi}} & z_j = n_1 + j. \end{cases}
 \end{aligned}$$

613 7.3 Summary of Fast Beta Linkage Method

```

1: procedure HASHING AND PREPROCESSING
2:   Construct and enumerate set of unique patterns  $\mathcal{P}$  from  $F$  and  $\{L_f\}$ .
3:   Partition files  $\mathbf{X}_1$  and  $\mathbf{X}_1$  into chunks  $\{\mathbf{X}_{1n}\}, \{\mathbf{X}_{2m}\}$ .
4:   for each  $n, m$  do
5:     Create comparison vectors between  $\mathbf{X}_{1n}$  and  $\mathbf{X}_{2m}$ .
6:     Hash records to  $\mathcal{R}_{nm}$  and calculate summary statistics  $\mathcal{H}_{nm}$ .
7:     Use SEI to reduce memory usage;  $\mathcal{R}_{nm} \rightarrow \mathcal{R}_{nm}^{\text{SEI}}$ .
8:   end for
9:   Synthesize results across pairings to get  $\tilde{\Gamma} = \{\mathcal{P}, \mathcal{R}, \mathcal{H}\}$ .
10: end procedure
11: procedure GIBBS SAMPLING
12:   Initialize  $\Phi$ , and  $Z$  parameters.
13:   for  $t \in \{1, \dots, T\}$  do
14:     Sample  $\Phi^{t+1} | Z^t, \tilde{\Gamma}$ .
15:     Sample  $h(Z^{t+1}) | \Phi^{t+1}, \tilde{\Gamma}$ . ▷ Sample agreement pattern, not record
16:   end for
17:   Sample  $Z | h(Z), \tilde{\Gamma}$ . ▷ Fills in record label based on agreement pattern
18: end procedure

```

7.4 One Hot Encoding Transformation

`fabl` makes use of one-hot encodings to aid in vectorized computations. For γ_{ij}^f with L_f levels, define e_{ij}^f to be an $L_f \times 1$ vector. When $\gamma_{ij}^f = \ell$, we set the ℓ^{th} element of e_{ij}^f to be 1, and set the other $L_f - 1$ elements of e_{ij}^f to be 0. We then concatenate the e_{ij}^f for all $f \in \{1, \dots, F\}$, resulting in the one-hot encoded comparison vector e_{ij} of length $\sum_{f=1}^F L_f$.

For example, consider comparing the toy records shown in Table 6 with $L = (3, 3, 2, 2)$ levels of agreement for last name, first name, DOB, and city respectively. Since the first name differs by only one letter, a reasonable comparison vector for this pair would be $\gamma_{ij} = (1, 2, 1, 2)$. The one hot encoding representation of this vector is $e_{ij} = (1, 0, 0, 0, 1, 0, 1, 0, 0, 1)$.

Last Name	First Name	DOB	City
Smith	Taylor	01/01/2000	Durham
Smith	Tayler	01/01/2000	Raleigh

Table 6: Example records for one hot encoding.

7.5 Traceplots for Simulation Study

Below are traceplots for one of the 900 linkage tasks that comprise the simulation in Section 5.2. It is set up with one error across the linkage fields and 50 duplicates across files. Traceplots across other settings exhibit similar behavior. Note that traceplots for `u` parameters show very little variation because the overwhelming majority of record pairs are nonmatching.




Figure 8: Representative traceplot of overlap between files from simulation study in Section 5.2




Figure 9: Representative traceplot of m parameter from simulation study in Section 5.2




Figure 10: Representative traceplot of u parameters from simulation study in Section 5.2

7.6 Additional Speed Simulation Study

To illustrate that different constructions of the comparison vectors lead to similar speed gains, we replicate the speed study of Section 5.1 under different settings. Here, we use four fields of comparison, each with three possible levels of agreement, resulting in $3^4 = 81$ possible patterns.




Figure 11: Run-time for BRL and fabl to run 1000 Gibbs iterations, including hashing step for fabl, for increasing values of both n_1 and n_2 . We see near quadratic growth in run-time for BRL, and near linear growth for fabl.




Figure 12: Run-time for BRL and fabl to run 1000 Gibbs iterations, including hashing step for fabl, with increasing n_1 , and n_2 fixed at 500. We see linear growth in run-time for BRL, and near constant run-time for fabl.

636 7.7 Traceplots for El Salvador Case Study




Figure 13: Traceplot for number of matches found across datasets in El Salvador case study




Figure 14: Traceplot for m parameter in El Salvador case study




Figure 15: Traceplot for u parameter in El Salvador case study