

Supplementary Material for “Efficient and Scalable Bipartite Matching with Fast Beta Linkage (fabl)”

Brian Kunderinger^{*}, Jerome P. Reiter^{*} and Rebecca C. Steorts[†]

Supplement A: Derivations of Full Conditionals

We provide detailed derivations of the full-conditionals provided in Section ???. In our derivations below, we assume that the new update to the Gibbs sampler is $s + 1$ and the current state of the sampler is s . As to not overload notation, we suppress the state of the sampler until presenting the final update of the full conditional that appears in Section ??.

The \mathbf{m} and \mathbf{u} parameters are updated through standard multinomial-Dirichlet distributions. For a particular m_{fl} , we have

$$\mathcal{L}(m_{fl}|\gamma, \alpha_{fl}, \mathbf{u}, \mathbf{Z}, \pi) \propto \prod_{i=1}^{n_A} \prod_{j=1}^{n_B} m_{fl}^{I(Z_j=i)I(\gamma_{ij}^f=l)I_{obs}(\gamma_{ij}^f)} m_{fl}^{\alpha_{fl}-1} = m_{fl}^{\alpha_{fl}(\mathbf{Z})-1},$$

where $\alpha_{fl}(\mathbf{Z}) = \alpha_{fl} + \sum_{i,j} I_{obs}(\gamma_{ij}^f)I(\gamma_{ij}^f = l)I(Z_j = i)$. Analogous procedures lead to the posterior distribution $\mathcal{L}(u_{fl}|\gamma, \beta_{fl}, \mathbf{m}, \mathbf{Z}, \pi) \propto u_{fl}^{\beta_{fl}(\mathbf{Z})-1}$, where $\beta_{fl}(\mathbf{Z}) = \beta_{fl} + \sum_{i,j} I_{obs}(\gamma_{ij}^f)I(\gamma_{ij}^f = l)I(Z_j \neq i)$. Thus for the vectors of parameters \mathbf{m}_f and \mathbf{u}_f , we have

$$\begin{aligned} \mathbf{m}_f^{(s+1)}|\mathbf{Z}^{(s)}, \gamma &\sim \text{Dirichlet}(\alpha_{f1}(\mathbf{Z}^{(s)}), \dots, \alpha_{fL_f}(\mathbf{Z}^{(s)})), \\ \mathbf{u}_f^{(s+1)}|\mathbf{Z}^{(s)}, \gamma &\sim \text{Dirichlet}(\beta_{f1}(\mathbf{Z}^{(s)}), \dots, \beta_{fL_f}(\mathbf{Z}^{(s)})). \end{aligned}$$

Since π encodes the rate of matching across the two data files, the full conditional $p(\pi|\gamma, \mathbf{Z}, \mathbf{m}, \mathbf{u}, \alpha_\pi, \beta_\pi)$ depends only on the number of links $n_{AB}(\mathbf{Z}) = \sum_{i=1}^{n_B} I(Z_j < n_A + j)$ encoded by \mathbf{Z} and hyperparameters. We have the full conditional

$$\begin{aligned} p(\pi|\mathbf{Z}, \alpha_\pi, \beta_\pi) &\propto p(\mathbf{Z}|\pi)p(\pi) \\ &\propto \pi^{n_{AB}(\mathbf{Z})}(1-\pi)^{n_B-n_{AB}(\mathbf{Z})}\pi^{\alpha_\pi-1}(1-\pi)^{\beta_\pi-1} \\ &\propto \pi^{n_{AB}(\mathbf{Z})+\alpha_\pi-1}(1-\pi)^{n_A-n_{AB}(\mathbf{Z})+\beta_\pi-1}. \end{aligned}$$

^{*}Department of Statistical Science, Duke University, P.O. Box 90251, Durham, NC 27708, USA brian.kunderinger@duke.edu, jreiter@duke.edu

[†]Departments of Statistical Science and Computer Science, Duke University, P.O. Box 90251, Durham, NC 27708, USA beka@stat.duke.edu

2 Efficient and Scalable Bipartite Matching with Fast Beta Linkage (fabl)

Thus, $\pi^{(s+1)} | \mathbf{Z}^{(s+1)}, \alpha_\pi, \beta_\pi$ has a $\text{Beta}(n_{AB}(\mathbf{Z}^{(s)}) + \alpha_\pi, n_B - n_{AB}(\mathbf{Z}^{(s)}) + \beta_\pi)$ distribution.

To express the full conditional for \mathbf{Z} , we consider the cases where Z_j does not have a link in A and where Z_j does have a link in A . Because we sample each Z_j independently of all other $Z_{j'}$ (for $j \neq j'$), we need only the full conditional for each individual Z_j . Let $\Gamma_{\cdot j}$ denote the matrix of n_A comparison vectors relating to record B_j . Following the observation of \mathbf{Z} , when B_j does not link to any record in A , the contribution to the likelihood is simply a product of u parameters, which we will call c_j :

$$p(\Gamma_{\cdot j} | \mathbf{m}, \mathbf{u}, \pi, Z_j = n_A + j) = \prod_{i=1}^{n_A} \prod_{f=1}^F \prod_{l=1}^{L_f} u_{fl}^{I(\gamma_{ij}^f=l)I_{obs}(\gamma_{ij}^f)} = c_j. \quad (\text{SI0})$$

When $Z_j = q$ for some $q \leq n_A$, we have

$$p(\Gamma_{\cdot j} | \mathbf{m}, \mathbf{u}, \pi, Z_j = q) = \prod_{f=1}^F \prod_{l=1}^{L_f} m_{fl}^{I(\gamma_{qj}^f=l)I_{obs}(\gamma_{qj}^f)} \prod_{i \neq q} \prod_{f=1}^F \prod_{l=1}^{L_f} u_{fl}^{I(\gamma_{ij}^f=l)I_{obs}(\gamma_{ij}^f)}. \quad (\text{SI0})$$

We multiply and divide by the u parameters for the matching record pair to obtain

$$\begin{aligned} p(\Gamma_{\cdot j} | \mathbf{m}, \mathbf{u}, \pi, Z_j = q) &\propto \prod_{f=1}^F \prod_{l=1}^{L_f} \left(\frac{m_{fl}}{u_{fl}} \right)^{I(\gamma_{qj}^f=l)I_{obs}(\gamma_{qj}^f)} \prod_{i=1}^{n_A} \prod_{f=1}^F \prod_{l=1}^{L_f} u_{fl}^{I(\gamma_{ij}^f=l)I_{obs}(\gamma_{ij}^f)} \\ &= w_{qj} c_j. \end{aligned} \quad (\text{SI0})$$

We can divide the result of each case by c_j to get

$$p(\Gamma_{\cdot j} | \mathbf{m}, \mathbf{u}, \pi, Z_j) \propto \begin{cases} w_{qj}, & q \leq n_A; \\ 1, & q = n_A + j. \end{cases} \quad (\text{SI0})$$

Lastly, we multiply the likelihood by the fast beta prior in (??) to obtain the full conditional

$$p\left(Z_j^{(s+1)} = q | \gamma, \mathbf{m}^{(s+1)}, \mathbf{u}^{(s+1)}, \pi^{(s+1)}\right) \propto \begin{cases} \frac{\pi^{(s+1)}}{n_A} w_{qj}^{(s+1)}, & q \leq n_A; \\ 1 - \pi^{(s+1)}, & q = n_A + j. \end{cases} \quad (\text{SI0})$$

Observe that the full conditional in (SI0) does not depend on $\mathbf{Z}^{(s)}$, implying that

$$p(\mathbf{Z}^{(s+1)} | \gamma, \mathbf{m}^{(s+1)}, \mathbf{u}^{(s+1)}, \pi) \propto \prod_q \left(\prod_j p\left(Z_j^{(s+1)} = q | \gamma, \mathbf{m}^{(s+1)}, \mathbf{u}^{(s+1)}, \pi\right) \right). \quad (\text{SI0})$$

Therefore, we update the full conditional for $\mathbf{Z}^{(s+1)}$ in parallel. This contrasts the Gibbs sampler of \mathbf{Z} , where the update to the matching matrix is sequential and accomplished using a partially collapsed Gibbs sampler, where one integrates out π . We do not opt for this approach as it would lead to similar sequential updates, which are computationally slow.

22 Supplement B: Bayes Estimate

We calculate a Bayes estimate \hat{Z} for the linkage parameter Z by assigning different positive losses to different types of errors, and minimizing posterior expected loss. We adopt the loss function proposed in ? in which $\hat{Z}_j \in \{1, \dots, n_A, n_A + j, R\}$, with R representing the option to leave the matching undetermined by the model. Specifically, we have

$$L(\hat{Z}_j, Z_j) = \begin{cases} 0, & \text{if } Z_j = \hat{Z}_j; \\ \theta_R, & \text{if } \hat{Z}_j = R; \\ \theta_{10}, & \text{if } Z_j \leq 1, \hat{Z}_j = n_A + j; \\ \theta_{01}, & \text{if } Z_j = n_A + j, \hat{Z}_j \leq n_A; \\ \theta_{11}, & \text{if } Z_j \leq n_A, \hat{Z}_j \leq n_A, Z_j \neq \hat{Z}_j. \end{cases}$$

23 Here, θ_R is the loss from not making a decision on the linkage status, θ_{10} is the loss
 24 from a false nonmatch, θ_{01} is the loss from a false match, and θ_{11} is the loss from the
 25 special case of a false match in which the record has a true match other than the one
 26 estimated by the model.

In general, we set $(\theta_{10}, \theta_{01}, \theta_{11}, \theta_R) = (1, 1, 2, \infty)$ inducing the decision rule

$$\hat{Z}_j = \begin{cases} i, & \text{if } p(Z_j = i | \gamma) > \frac{1}{2}; \\ 0, & \text{otherwise.} \end{cases}$$

27 Since **fab1** does not strictly enforce one-to-one matching, it is possible for this Bayes
 28 estimate to link multiple records in B to one record in A . In the event that we have two
 29 records B_j and $B_{j'}$ such that both $p(\hat{Z}_j = i | \gamma) > \frac{1}{2}$ and $p(\hat{Z}_{j'} = i | \gamma) > \frac{1}{2}$, we accept the
 30 match with the higher posterior probability, and declare the other to have no match.
 31 Since each Z_j is independent, this is equivalent to minimizing the expected loss subject
 32 to the constraint that $\hat{Z}_j \neq \hat{Z}_{j'}$ for all $j \neq j'$. A similar approach appears in the most
 33 probable maximal matching sets used by ? to match records to latent entities.

34 When we seek a partial estimate of the linkage structure, leaving a portion of record
 35 pairs to be classified manually in clerical review, we adopt losses $(\theta_{10}, \theta_{01}, \theta_{11}, \theta_R) =$
 36 $(1, 1, 2, .1)$. For a more in-depth explanation of this function and the induced Bayes
 37 estimate, see ?.

38 Supplement C: Accuracy under Partial Estimates

39 In this section, we repeat the simulation study in Section ?? of the main text, allowing
 40 for clerical review rather than forcing all records to have or not have links. Specifically,
 41 by leaving $\theta_{10} = \theta_{01} = 1$ and $\theta_{11} = 2$, but setting $\theta_R = 0.1$, we allow the model
 42 to decline to decide a match for certain records, with nonassignment being 10% as
 43 costly as a false match. In this context, we are no longer focused on finding all true
 44 matches, but rather protecting against false matches. Thus, instead of recall, we use the
 45 negative predictive value (NPV), defined as the proportion of non-links that are actual
 46 nonmatches. Mathematically, $\text{NPV} = \sum_{j=1}^{n_B} I(\hat{Z}_j = Z_j = n_A + j) / \sum_{j=1}^{n_B} I(\hat{Z}_j = n_A + j)$.

4 Efficient and Scalable Bipartite Matching with Fast Beta Linkage (*fabl*)




Figure SI1: Negative predictive value (NPV), positive predictive value (PPV), and decision rate (DR) on data files in the simulation in Appendix C. We see poorer performance for **fabl** only in situations with high overlap.

We continue to use the precision, which is renamed the positive predictive value (PPV) in this context. Lastly, we also examine the rejection rate (RR), or how often the model declines to make a linkage decision, defined as $RR = \sum_{j=1}^{n_B} I(\hat{Z}_j = R)/n_B$. To convey this information alongside NPV and PPV, for which values close to 1 indicate strong performance, we report the decision rate (DR), defined as $DR = 1 - RR$.

In Figure SI1, we see that **fabl** maintains equivalently strong PPV as **BRL** across all linkage settings. However, with high amounts of error, and thus fewer accurate and discerning fields of information, the rejection rate under **fabl** rises, leading to a decrease in NPV. Since **fabl** does not remove previously matched records from consideration for a new record, posterior probabilities of matches at times can be split across more records; in contrast, **BRL** is able to maintain higher confidence in matches in this setting. If one wishes to use partial estimates, **fabl** will possibly leave more linkages for the modeler to match by hand than would be left under **BRL**, but the decisions made by each method should have nearly equal accuracy.

Supplement D: Traceplots for Simulation Study

Figures SI2, SI3, and SI4 are traceplots for one of the 900 linkage tasks that comprise the simulation in Section ?? . It is set up with one error across the linkage fields and 50

64 duplicates across files. Traceplots across other settings exhibit similar behavior. Note
65 that traceplots for u parameters show very little variation because the overwhelming
66 majority of record pairs are nonmatching.




Figure SI2: Representative traceplot of overlap between files from simulation study in Section ??.

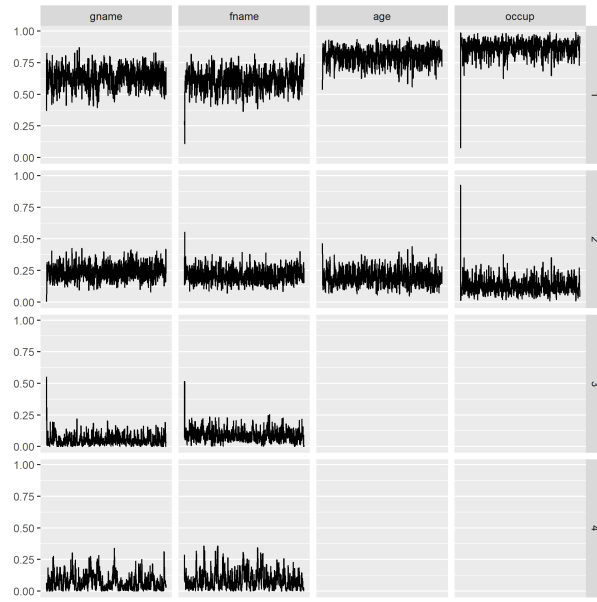


Figure SI3: Representative traceplot of m parameter from simulation study in Section ??.




Figure SI4: Representative traceplot of \mathbf{u} parameters from simulation study in Section ??.

	\mathbf{m}			\mathbf{u}		
	Agree	Partial	Disagree	Agree	Partial	Disagree
Feature 1	$\frac{9}{10}$	$\frac{9}{100}$	$\frac{1}{100}$	$\frac{1}{100}$	$\frac{3}{100}$	$\frac{96}{100}$
Feature 2	$\frac{9}{10}$	$\frac{9}{100}$	$\frac{1}{100}$	$\frac{1}{100}$	$\frac{3}{100}$	$\frac{96}{100}$
Feature 3	$\frac{9}{10}$	$\frac{9}{100}$	$\frac{1}{100}$	$\frac{1}{100}$	$\frac{3}{100}$	$\frac{96}{100}$
Feature 4	$\frac{9}{10}$	$\frac{9}{100}$	$\frac{1}{100}$	$\frac{1}{100}$	$\frac{3}{100}$	$\frac{96}{100}$

Table SI1: Probabilities used for \mathbf{m} and \mathbf{u} distributions in simulation study in Appendix E.

Supplement E: Additional Speed Simulation Study

Figures SI5 and SI6 illustrate that different constructions of the comparison vectors lead to similar speed gains. We replicate the speed study of Section ?? under different settings. Here, we use four fields of comparison, each with three possible levels of agreement, resulting in $3^4 = 81$ possible patterns. The \mathbf{m} and \mathbf{u} parameters for this simulation are shown Table SI1.




Figure SI5: Run-time for BRL and `fabl` to run 1000 Gibbs iterations in simulation study in Appendix E, including hashing step for `fabl`, for increasing values of both n_A and n_B . We see near quadratic growth in run-time for BRL, and near linear growth for `fabl`.




Figure SI6: Run-time for BRL and `fabl` to run 1000 Gibbs iterations in simulation study in Appendix E, including hashing step for `fabl`, with increasing n_A and n_B fixed at 500. We see linear growth in run-time for BRL, and near constant run-time for `fabl`.

73 **Supplement F: Traceplots for El Salvador Case Study**






Figure SI7: Traceplot for number of matches found across data files in El Salvador case study.

Figure SI8: Traceplot for m parameter in El Salvador case study.Figure SI9: Traceplot for u parameter in El Salvador case study.