Efficient and Scalable Bipartite Matching with Fast Beta Linkage (fabl)

Brian Kundinger[a]   Jerome Reiter[a]   Rebecca C. Steorts [b]

[a] Department of Statistical Science, Duke University
[b] Department of Statistical Science and Computer Science, Duke University
Principal Mathematical Statistician, United States Census Bureau

October 9, 2021

Abstract

Thanks Beka for writing that first abstract. I think this one captures the idea better, even if its not exactly in "outline" form.

In this paper, we propose *fast beta linkage* (`fabl`), which extends a recent Bayesian Fellegi Sunter model for increased efficiency and scalability. Specifically, we relax the one-to-one matching requirement of the Beta Record Linkage method of (Sadinle, 2017) and propose independent priors over the matching space. This allows us to employ hashing techniques that hasten calculations and reduce computational complexity, complete pairwise record comparisons over large datasets through parallel computing, and reduce memory costs through a new technique called *storage efficient indexing*. Through simulation studies and a case study of homicides from the El Salvadoran Civil War, we show that our method has markedly increased speed with minimal loss of accuracy.

Keywords: record linkage, bipartite record linkage, Bayesian methods, Gibbs sampling, hashing techniques, Markov chain Monte carlo, parallel/distributed computing

1 INTRODUCTION

I have some proposed revisions for the first part of the paper. In the introduction, I reordered things and changed some topic sentences to highlight that the main contribution here is speed and scale.

Also, in the prior version, we didn't get to any of the proposed material until Section 4.3, which did not seem reasonable to me. I know you recommended splitting the review of the literature between the Introduction and Section 2 on bipartite matching, but since we needed the mathemetical overview of FS and Sadinle, I decided to move all of the literature review into the introduction. I've seen this done in many papers.

Almost all of the content is still there, but it's moved around in a way that makes more sense to me. Let me know your thoughts! (But please avoid hard edits until we talk on Monday!)

I have some additional edits in the later part of the paper about the partial Bayes estimate, computational complexity, and added detail for the case study. All of that is in blue as well.

————————————————————

Record linkage is the task of identifying duplicate records across multiple data sources, often in the absence of a unique identifier (Christen, 2012). Record linkage is an increasingly important task in "data cleaning," that is used for inferential and predictive analyses in fields such as statistics, computer science, machine learning, political science, economics, precision medicine, official statistics, and others. In this paper, we consider bipartite record linkage, which is the merging two databases that contain duplication across the two databases, but do not contain any duplications within each database.

Many statistical record linkage methods and case studies are extensions of seminal work of Fellegi and Sunter (1969) and Newcombe et al. (1959). Specifically, Fellegi and Sunter (1969) created comparison vectors for each pair of records in the data and independently classified those pairs as a match or a non-match using a likelihood ratio test (or hypothesis test). Recent work in the statistical literature has extended this approach for a wide variety of applications (Winkler and Thibaudeau, 1991; Fair, 2004; Wagner et al., 2014; Gill and

Goldacre, 2003). Although the need to compute comparisons for all record pairs initially restricted this family of methods, recent work by Enamorado et al. (2019a) used innovative hashing techniques to scale the approach to handle large datasets.

The independent pairwise matching assumption from Fellegi and Sunter is popular mainly for its mathematical simplicity, but is often unreasonable in practice. in many contexts, we know that there are no duplications within a database, meaning that one record from one database should be linked with at most one record in the other database. Here, additional declared matches are known by assumption to be false. Many extensions to Fellegi and Sunter (1969) resolve these false matches as a post-processing step (Jaro, 1989; Enamorado et al., 2019a), but this model misspecification can still lead to poor results Sadinle (2017).

Alternatively, one can embed one-to-one matching requirements into the model specification itself, at an additional computational cost. Larsen (2005a) employed a Metropolis-Hastings algorithm to only allow sampling from matches that respected one-to-one matching, but such algorithms are slow due to the combinatorial nature of the constrained matching space. Fortunato (2010) used simulated annealing to target the space of permitted matches under this constraint, but the method was so computationally intensive it could used to link databases of less than 100 records. Most recently, Sadinle (2017) proposed *Beta Record Linkage* (BRL), using an innovative prior over the space of bipartite matchings to strictly enforce one-to-one reqirements throughout his Gibbs sampler. Additionally, he introduces a class of loss functions that allows for a flexible estimation of the linkage structure, such that the modeller can weight the relative importance of false positives and false negatives, and identify records pairings to be be decided through clerical review. Unfortunately, the method becomes slow when working with larger datasets, and unusable when the number of pairwise record comparisons becomes too large.

In this paper, we propose *fast beta linkage* (`fabl`), which extends a `BRL` model for increased efficiency and scalability. Specifically, we relax the one-to-one matching requirement of the Beta Record Linkage method of (Sadinle, 2017) and propose independent priors over the matching space. This allows us to employ hashing techniques that hasten calculations

3

and reduce computational complexity, complete pairwise record comparisons over large datasets through parallel computing, and reduce memory costs through a new technique called *storage efficient indexing.* Throughout, we argue that even in cases where a bipartite matching is desired, this relaxed model (with a simple post-processing procedure) provides equivalently accurate estimation of the linkage structure, more information through which to asses model misspecification, and greatly enhanced speed. To recognize the lineage from the original `BRL` method, we name our method *fast beta linkage* (`fabl`, pronounced "fable")

In what follows, we first review more thoroughly the Fellegi Sunter and Beta Record Linkage methods. Then, we provide the model specification under `fabl` and derive the Gibbs sampler for posterior inference. Next, we introduce the hashing technique and storage efficient indexing used to increase the speed of calculations and the scale of the linkage tasks we can undertake. Lastly, we use simulation studies and a case study of homicides from the El Salvadoran Civil War, we show that our method has markedly increased speed with minimal loss of accuracy when compared to `BRL`.

2   REVIEW OF PRIOR WORK

Consider two databases $\boldsymbol{X}_1$ and $\boldsymbol{X}_2$ with respective sizes n_1 and n_2. Without loss of generality, denote files suchs that $n_1 \geq n_2$. In the context of bipartite matching, we assume that there are duplications across the databases but not within each database. Within this framework, the set of matches across datasets can be viewed in two equivalent ways:

$$\Delta_{ij} = \begin{cases} 1 & \text{if records } i \in \boldsymbol{X}_1 \text{ and } j \in \boldsymbol{X}_2 \text{ refer to the same entity;} \\ 0 & \text{otherwise.} \end{cases} \tag{1}$$

This matrix valued representation of the set of matches is of size $n_1 \times n_2$, which can become costly for large linkage tasks or for long posterior samplers. More compactly, bipartite matching can also be viewed as a labeling $\boldsymbol{Z} = (Z_1, \ldots, Z_{n_2})$ for the records in database $\boldsymbol{X}_2$ such that

$$Z_j = \begin{cases} i & \text{if records } i \in \boldsymbol{X}_1 \text{ and } j \in \boldsymbol{X}_2 \text{ refer to the same entity;} \\ n_1 + j & \text{if records } j \in \boldsymbol{X}_2 \text{ does not have a match in database } \boldsymbol{X}_1. \end{cases} \tag{2}$$

4

Depending on which representation is the most convenient, we can go back and forth between the two using $\Delta_{ij} = I(Z_j = i)$, where $I(\cdot)$ is the indicator function.

Denote the set of matches by $\boldsymbol{M} = \{(i, j) : i \in \boldsymbol{X}_1, j \in \boldsymbol{X}_2, \Delta_{ij} = 1\}$. Denote the set of non-matches $\boldsymbol{U} = \{(i, j) : i \in \boldsymbol{X}_1, j \in \boldsymbol{X}_2, \Delta_{ij} = 0\}$. Merging two databases and removing duplications can be viewed as identifying the sets of matches/non-matches, or rather $\boldsymbol{M}, \boldsymbol{U}$. We refer to record pairs that are estimated as matches as "links", and record pairs that are estimated as non-matches are known as "non-links". Fellegi and Sunter (1969) proposed using all-to-all pairwise comparison of records to estimate the matches and non-matches.

2.1 Comparison Data

Intuitively, records that are co-referent or rather refer to the same entity should be similar; records that are non co-referent and refer to different entities should not be similar. A common way to encode this is using a comparison vector $\boldsymbol{\gamma}_{ij}$ which is computed for each record pair (i, j) in $\boldsymbol{X}_1 \times \boldsymbol{X}_2$. Denote F as the criteria for compare the records such that $\boldsymbol{\gamma}_{ij} = (\gamma_{ij}^1, \gamma_{ij}^2, \ldots, \gamma_{ij}^f, \ldots \gamma_{ij}^F)$. In most cases, F represents to a single comparison per feature that the databases share, such as gender.

The simplest way to compare two records is check for agreement or disagreement, and this is commonly used for nominal variables such as gender. For more complex measurements, such as textual or numeric feature information, we can take into account partial agreement patterns using common string metrics, such as the Edit, Jaro, or Jaro-Winkler distance functions. This allows us to divide our records into a set of similarity values into different levels of disagreement (Bilenko and Mooney, 2006; Elmagarmid et al., 2007).

Let $\mathcal{S}_f(i, j)$ denote a general similarity measure for feature f of records i and j, where the range of $\mathcal{S}_f$ can be divided into $L_f + 1$ intervals denoted by $I_{f0}, I_{f1}, \ldots, I_{fL_f}$, representing disagreement levels. Following convention, I_{f0} represents the highest level of agreement (inclusive of complete agreement) and I_{fL_f} represents the highest level of disagreement (including complete disagreement). Thus, we can construct comparison vectors in the following way:

$$\gamma_{ij}^f = \ell \text{ if } \mathcal{S}_f(i, j) \in I_{f\ell}.$$

5

The choice of $I_{f\ell}$ are application specific, which we discuss in our simulation and case study.

2.2 Fellegi-Sunter Model

While we have introduced comparison data, $\boldsymbol{\gamma}_{ij}$, this is not sufficient to determine whether a record is a match/non-match given errors that naturally occur in the data. This motivated Fellegi and Sunter (1969) to consider the following likelihood ratio

$$w_{ij} = \frac{P(\gamma_{ij} \mid \Delta_{ij} = 1)}{P(\gamma_{ij} \mid \Delta_{ij} = 0)} \tag{3}$$

as a weight to estimate if record pairs are a match/non-match. The ratio will be large if there is strong evidence of the pair being a match, and small otherwise.

In practice $P(\cdot \mid \Delta_{ij} = 1)$ and $P(\cdot \mid \Delta_{ij} = 0)$ are unknown to the modeler, and thus must be estimated. We assume the comparison vectors are independent given the bipartite matching and assume that the match status of the record pairs are independent. It then follows that we can model the comparison data using mixture models in the following way:

$$\begin{aligned} \Gamma_{ij} = \gamma_{ij} \mid \Delta_{ij} = 1 &\overset{iid}{\sim} \mathcal{M}(\boldsymbol{m}), \\ \Gamma_{ij} = \gamma_{ij} \mid \Delta_{ij} = 0 &\overset{iid}{\sim} \mathcal{U}(\boldsymbol{u}), \\ \Delta_{ij} &\overset{iid}{\sim} \text{Bernoulli}(p), \end{aligned} \tag{4}$$

where p is the proportion of records that match. Note that $\boldsymbol{m}$ and $\boldsymbol{u}$ are vectors of the match/non-match parameters. One common way to estimate the unknown parameters of the aforementioned model is using the EM algorithm.

After estimating $\boldsymbol{m}$ and $\boldsymbol{u}$, the Fellegi Sunter method uses thresholds T_m and T_u such that all record pairs with $w_{ij} > T_m$ are declared as links, and all those with $w_{ij} < T_u$ declares as nonlinks (with those such that $T_u < w_{ij} < T_m$ left undetermined). This is a problem in practice as transitive closures are violated. Since each w_{ij} is calculated independently, there are often situation in which both $w_{ij} > T_m$ and $w_{i'j} T_m$, even those such a matching is not allowed by assumption.

Jaro (1989) proposed addressing this issue via a post-processing optimization approach as follows:

$$\max_{\Delta} \sum_{i=1}^{n_1} \sum_{j=1}^{n_2} w_{ij} \Delta_{ij} \quad \text{subject to} \quad \Delta_{ij} \in \{0, 1\}; \tag{5}$$

$$\sum_{i=1}^{n_1} \Delta_{ij} \leq 1, j = 1, 2, \dots n_2; \text{ and}$$

$$\sum_{i=1}^{n_2} \Delta_{ij} \leq 1, j = 1, 2, \dots n_1.$$

Equation 10 ensures that each record of $\boldsymbol{X}_2$ only matches with at most one record in $\boldsymbol{X}_1$ (and the reverse condition). This is commonly known as the *maximum-weight bipartite matching problem* or a *linear sum assignment problem*. The output of this algorithm is a bipartite matching that maximizes the sum of the weights among matched pairs. Jaro (1989) provided no theoretical justification for using this approach, however, Sadinle (2017) recently showed that under certain conditions, equation 10 is the maximum likelihood estimate.

2.3   Beta Linkage Model

One criticism of using mixture models is the fact that the decision rule will lead to "many-to-many assignments", breaking transitive closures. One way of resolving this issue to incorporate the constraint directly into the model. The Beta Linkage Model of Sadinle (2017) does through the following model:

$$\boldsymbol{\Gamma}_{ij} \mid Z_j = i \overset{iid}{\sim} \mathcal{M}(\boldsymbol{m}) \tag{6}$$

$$\boldsymbol{\Gamma}_{ij} \mid Z_j \neq i \overset{iid}{\sim} \mathcal{U}(\boldsymbol{u})$$

$$\boldsymbol{Z} \sim \mathcal{B},$$

where $\mathcal{B}$ represents a prior on the space of bipartite matchings and $\mathcal{M}(\boldsymbol{m}), \mathcal{M}(\boldsymbol{u})$ are models for the comparison vectors of matches and non-matches.

Here, $\mathcal{B}$ is a prior distribution on the entire $\mathbf{Z}$ vector, called the *beta distribution for bipartite matching*, given by

$$P(\mathbf{Z}|\alpha_\pi, \beta_\pi) = \frac{(n_1 - n_{12}(\mathbf{Z}))!}{n_1!} \frac{\mathrm{B}(n_{12}(\mathbf{Z}) + \alpha_\pi, n_2 - \mathbf{Z}) - \beta_\pi])}{\mathrm{B}(\alpha_\pi, \beta_\pi)}$$

where $B(\cdot, \cdot)$ represents the Beta function. This prior induces a Gibbs sampler that strictly enforces one-to-one matching, previously matched records from the set of candidate records when sampling Z_j. This creates a dependency that makes the sampler *inherently serial*, which makes the sampler slow when working on large linkage tasks.

2.4   Blocking

In practice, it is not feasible to make all-to-all record comparisons as the computationally complexity is of the order $O(n_1 \times n_2)$. The most common solution is utilize blocking, which places similar records into partitions, or "blocks" to reduces this computational burden. In deterministic blocking, the modeller chooses a field thought to be highly reliable, and only compares records that agree on that field. The record linkage method is then applied independently across all blocks, which can be done in parallel for additional speed gains. However, blocking on an unreliable field can lead to many missed matches, making this form of blocking undesirable in many situations.

In this paper, we use hashing and new technique called storage efficient indexing to increase the scale of the linkage tasks we can undertake without the need for blocking. This is useful when there is no reliable blocking field available, or one desires estimates of model parameters for the entire sample in question. In practice, the method the follows can be combined with blocking for even greater computational gains, but all derivations, simulations, and case studies are presented without blocking.

———————————————————————

I think its best just not to mention probabalistic blocking. My method doesn't use it, and introducing it is really confusing (Sadinle 2017, Enamorado 2019, and Steorts 2014/2015 don't mention it).

We note here that the word "hashing" is also used in probabalistic blocking, which uses techniques from machine learning to cluster records together before attempting to link records. The use of such probabalistic blocking to provide even greater computational gains will be explored in future work.

———————————————————————

3 FAST BETA LINKAGE

Recall our compact representation of the matching structure $\boldsymbol{Z} = (Z_1, Z_2, \ldots Z_{n_2})$ with $Z_j \in \{1, 2, 3, n_1, n_1 + j\}$, where $Z_j < n_1$ indicates a link and $Z_j = n_1 + j$ indicates a nonlink. In constrast to the prior over the vector $\boldsymbol{Z}$ from Sadinle (2017), we propose independent priors for each component Z_j. We denote the *fast beta prior* as follows:

$$Z_j | \lambda = \begin{cases} \frac{1}{n_1} \lambda & z_j \leq n_1; \\ 1 - \lambda & z_j = n_1 + j \end{cases}$$

$$\lambda \sim \text{Beta}(\alpha_\lambda, \beta_\lambda)$$

Intuitively, this set of priors says that record $j \in \boldsymbol{X}_2$ has some match in $\boldsymbol{X}_1$ with probability λ, and that each record $i \in \boldsymbol{X}_1$ is equally likely to be that match. λ itself is given a prior distribution to reflect prior beliefs about the overall rate of matchingness in the data. One choice of uninformative prior might be $\lambda \sim \text{Beta}(1, 1)$, which corresponds to a prior belief that nonmatches and matches are equally likely, and another might be $\lambda \sim \text{Beta}\left(1, \frac{1}{n_1}\right)$, which corresponds to a uniform prior on the labeling of $\mathbf{Z}$.

Importantly, these independent priors relax the one-to-one requirement from BRL; our sampler does ensure that each record in $\boldsymbol{X}_2$ can be matched to at most one record in $\boldsymbol{X}_1$, but allows for the possibility that multiple records in $\boldsymbol{X}_2$ match to the same record in $\boldsymbol{X}_1$. This behavior may be desireable in its own right; Newcombe and Rhynas (1962) was interested in matching birth records to marriage certificates, a situation in which such "many-to-one" matchings are reasonable. In situations when a bipartite estimate is desired, we can use a simple post processing procedure shown in 6.5 to resolve illegal matchings.

3.1 Model Specification

NOTE: I really want to present the entire model in one place like this. I saw this for the first time in Steorts 2014, and I really liked it. Part of the reason it took me so long to understand Sadinle 2017 was because the model was presented over several sections

As commonly done in the record linkage literature, we assume that agreement levels across features are conditionally independent. Additionally, we assume that any missing data is missing at random. Let $m_{f\ell} = P(\Gamma^f_{ij} = \ell \mid Z_j = i)$ denote the probability of that matching records have level ℓ of disagreement in feature f, and $u_{f\ell} = P(\Gamma^f_{ij} = \ell \mid Z_j \neq i)$ represent the same probability for non-matches. In addition, denote $\boldsymbol{m}_f = (m_{f1}, \ldots, m_{fL_f})$, $\boldsymbol{u}_f = (u_{f1}, \ldots, u_{fL_f})$, $\boldsymbol{m} = (m_1, \ldots m_F)$, $\boldsymbol{u} = (u_1, \ldots u_F)$, and $\Phi = (\boldsymbol{m}, \boldsymbol{u})$. The model can then be specified as follows:

$$\mathcal{L}(\boldsymbol{Z}, \Phi \mid \boldsymbol{\Gamma}) = \prod_{j=1}^{n_1}\prod_{i=1}^{n_2}\left[\prod_{f=1}^{F}\prod_{l=1}^{L_f} m_{fl}^{I(Z_j=i)} u_{fl}^{I(Z_j\neq i)}\right]^{I(\gamma^f_{ij}=l)I_{obs}(\gamma^f_{ij})}$$

$$\boldsymbol{m}_f \sim \text{Dirichlet}(\alpha_{f0}, \ldots, \alpha_{fL_f})$$

$$\boldsymbol{u}_f \sim \text{Dirichlet}(\beta_{f0}, \ldots, \beta_{fL_f}).$$

$$Z_j|\lambda = \begin{cases} \frac{1}{n_1}\lambda & z_j \leq n_1; \\ 1-\lambda & z_j = n_1 + j \end{cases}$$

$$\lambda \sim \text{Beta}(\alpha_\lambda, \beta_\lambda)$$

3.2   Gibbs Sampler

We work with the following factorization of the joint posterior distribution:

$$p(\mathbf{Z}, \Phi, \lambda|\Gamma) \propto \mathcal{L}(\boldsymbol{Z}, \Phi \mid \Gamma^{obs})p(\mathbf{Z}|\lambda)p(\Phi)p(\lambda)$$

$$\propto \prod_{j=1}^{n_2}\prod_{i=1}^{n_1}\left[\prod_{f=1}^{F}\prod_{l=1}^{L_f} m_{fl}^{I(Z_j=i)} u_{fl}^{I(Z_j\neq i)}\right]^{I(\gamma^f_{ij}=l)I_{obs}(\gamma^f_{ij})}$$

$$\times \prod_{f=1}^{F}\prod_{l=1}^{L_f} m_{fl}^{\alpha_{fl}-1} \times \prod_{f=1}^{F}\prod_{l=1}^{L_f} u_{fl}^{\beta_{fl}-1}$$

$$\times \prod_{j=1}^{n_2}\left[I(Z_j \leq n_1)\frac{1}{n_1}\lambda + I(Z_j > n_1)(1-\lambda)\right]$$

This factorization leads to following Gibbs Sampler:

Sample $\Phi^{(s+1)}|\Gamma, \mathbf{Z}^{(s)}$: The $\mathbf{m}$ and $\mathbf{u}$ parameters are updated through standard multinomial-dirichlet mechanics. Thus we have

$$\mathbf{m}_f|\mathbf{Z}, \Gamma \sim \text{Dirichlet}(\alpha_{f1}(\mathbf{Z}), \ldots, \alpha_{fL_f}(\mathbf{Z}))$$

$$\mathbf{u}_f|\mathbf{Z}, \Gamma \sim \text{Dirichlet}(\beta_{f1}(\mathbf{Z}), \ldots, \beta_{fL_f}(\mathbf{Z}))$$

where $\alpha_{fl}(\mathbf{Z}) = \alpha_{fl} + \sum_{i,j} I_{obs}(\gamma_{ij}^f)I(\gamma_{ij}^f = l)I(z_j = i)$ and $\beta_{fl}(\mathbf{Z}) = \beta_{fl} + \sum_{i,j} I_{obs}(\gamma_{ij}^f)I(\gamma_{ij}^f = l)I(z_j \neq i)$.

Sample $\lambda^{(s+1)}|\mathbf{Z}^{(s)}$: As a function of λ, the linkage structure parameter $\mathbf{Z}$ is sequence of successes (when $z_j < n_A + 1$) and failures (when $z_j = n_A + 1$), and therefore $p(\mathbf{Z}|\lambda) = \mathcal{L}(\lambda|\mathbf{Z})$ is determined only by the number of duplicates $n_{12}(\mathbf{Z}) = \sum_{i=1}^{n_2} I(z_j < n_1 + 1)$ encoded by $\mathbf{Z}$. Thus we have

$$p(\lambda|\mathbf{Z}) \propto p(\mathbf{Z}|\lambda)p(\lambda)$$

$$\propto \lambda^{n_{12}(\mathbf{Z})}(1 - \lambda)^{n_2 - (n_{12}(\mathbf{Z})}\lambda^{\alpha_\lambda - 1}(1 - \lambda)^{\beta_\lambda - 1}$$

$$\propto \lambda^{n_{12}(\mathbf{Z}) + \alpha_\lambda - 1}(1 - \lambda)^{n_1 - n_{12}(\mathbf{Z}) + \beta_\lambda - 1}$$

$$\implies \lambda^{(s+1)}|\mathbf{Z}^{(s+1)} \sim \text{Beta}(D + \alpha_\lambda, n_2 - n_{12}(\mathbf{Z}) + \beta_\lambda)$$

Sample $\mathbf{Z}^{(s+1)}|\Gamma, \Phi^{(s+1)}, \lambda^{(s+1)}$: Because we sample Z_j independently of all other $Z_{j'}$, we use only the full conditional for an individual Z_j. Let $\Gamma_{.j}$ denote the set of n_A comparison vectors with $j \in B$, and note that as a function of Z_j, the likelihood $p(\Gamma_{.j}|Z_j, \mathbf{m}, \mathbf{u}) = \mathcal{L}(Z_j|\Gamma_{.j}, \mathbf{m}, \mathbf{u})$ is a discrete distribution with probabilities proportional to

$$p(\Gamma_{.j}|Z_j = z_j, \Phi) \propto \prod_{i=1}^{n_1} \left[\prod_{f=1}^{F} \prod_{l=1}^{L_f} m_{fl}^{I(Z_j=i)} u_{fl}^{I(Z_j \neq i)} \right]^{I(\gamma_{ij}^f=l)I_{obs}(\gamma_{ij}^f)}$$

$$\propto \prod_{i=1}^{n_1} \left(\prod_{f=1}^{F} \prod_{l=1}^{L_f} \frac{m_{fl}}{u_{fl}} \right)^{I(z_j=i)I(\gamma_{ij}^f=l)I_{obs}(\gamma_{ij}^f)}$$

$$= \begin{cases} w_{ij} & z_j \leq n_1; \\ 1 & z_j = n_1 + j \end{cases}$$

where $w_{ij} = \left(\frac{\prod_{f=1}^{F} \prod_{l=1}^{L_f} m_{fl}}{\prod_{f=1}^{F} \prod_{l=1}^{L_f} u_{fl}} \right)^{I(\gamma_{ij}^f=l)} = \frac{P(\gamma_{ij}|Z_j=i)}{P(\gamma_{ij}|Z_j \neq i)I_{obs}(\gamma_{ij}^f)}$. The interested reader should note that these are precisely the likelihood ratios used in the Fellegi-Sunter model to classify matches and non-matches, and we therefore refer to w_{ij} as the *Fellegi Sunter weights*.

With the likelihood in this form, we can derive the full conditional

$$p(Z_j|\Gamma_{.j}, \Phi, \lambda) \propto p(\Gamma_{.j}|Z_j, \Phi)P(Z_j|\lambda)$$

$$\propto \left(\sum_{i=1}^{n_A} w_{ij} \mathbf{1}_{z_j=i} + \mathbf{1}_{z_j=n_A+1} \right) \left(\lambda \sum_{i=1}^{n_A} \frac{1}{n_A} \mathbf{1}_{z_j=i} + (1-\lambda)\mathbf{1}_{z_j=n_A+1} \right)$$

$$= \frac{\lambda}{n_A} \sum_{i=1}^{n_A} w_{ij} \mathbf{1}_{z_j=i} + (1-\lambda)\mathbf{1}_{z_j=n_A+1}$$

$$\implies Z_j^{(s+1)}|\Phi, \Gamma, \lambda \propto \begin{cases} \frac{\lambda}{n_A} w_{ij} & z_j \leq n_A; \\ 1-\lambda & z_j = n_A + 1 \end{cases}$$

In order to make fair comparisons against the (Sadinle, 2017) model, we integrate over the posterior of λ and rearrange terms to produce the final full conditional:

$$p\left(Z_j^{(s+1)} = i | \mathbf{m}, \mathbf{u}, \mathbf{Z^{(s)}} \right) \propto \begin{cases} w_{ij} & i \leq n_A; \\ n_A \frac{n_B-D+\beta_\lambda}{D+\alpha_\lambda} & i = n_A + 1 \end{cases}$$

3.3 Bayes Estimate

We calculate a Bayes estimate $\hat{\mathbf{Z}}$ for the linkage parameter $\mathbf{Z}$ by assigning different positive losses to different types of errors, and minimizing posterior expected loss. We adopt the loss function proposed in Sadinle 2017, in which $\hat{Z}_j \in \{1, \ldots, n_A + 1, R\}$, with R representing the option to leave the matching undetermined by the model. Specifically, we have

$$L(\hat{Z}_j, Z_j) = \begin{cases} 0 & \text{if } Z_j = \hat{Z}_j; \\ \theta_R, & \text{if } \hat{Z}_j = R; \\ \theta_{10}, & \text{if } Z_j \leq n_A, \hat{Z}_j = n_A + 1; \\ \theta_{01}, & \text{if } Z_j = n_A + 1, \hat{Z}_j \leq n_A; \\ \theta_{11}, & \text{if } Z_j, \hat{Z}_j \leq n_A, Z_j \neq \hat{Z}_j; \end{cases}$$

Here, θ_R is the loss from not making a decision on the linkage status, θ_{10} is the loss from a false non-match, θ_{01} is the loss from a false match, and θ_{11} is the loss from the special case of a false match in which the record has a true match other than the one estimated by the model.

We seek to minimize the posterior expected loss, given by $\mathbb{E}[L(\hat{\mathbf{Z}}, \mathbf{Z}) | \Gamma^{obs}] = \sum_{\mathbf{Z}} L(\hat{Z}_j, Z_j) P(\mathbf{Z} | \Gamma^{obs})$. In this paper, we adopt losses $\theta_R = \infty, \theta_{10} = 1, \theta_{01} = 1, \theta_{11} = 2$, inducing the intuitive decision rule

$$\hat{Z}_j = \begin{cases} i, & \text{if } P(Z_j = i | \Gamma) > \frac{1}{2}; \\ 0, & \text{otherwise}; \end{cases}$$

For a more in-depth explanation of this function and the induced Bayes estimate, see (Sadinle, 2017).

Since our Gibbs procedure does not strictly enforce one-to-one matching, it is possible for this Bayes estimate to link multiple records in $\mathbf{X_2}$ to one record in $\mathbf{X_1}$. It is often desirable to see such matches, because they provide evidence of model misspecification; specifically, large numbers of "many to one" matches indicate that the assumption of no duplicates within files may not be reasonable. We will explore this issue more deeply through the El

Salvador homicide case study.

To achieve a Bayes estimate that fulfills one-to-one matching requirement, we simply minimize the expected loss subject to the constraint that $\hat{Z}_j \neq \hat{Z}_{j'}$ for all $j \neq j'$. In the event that two matches. In the event that we have two records j and j' such that both $P(Z_j = i|\Gamma) > \frac{1}{2}$ and $P(Z_{j'} = i|\Gamma) > \frac{1}{2}$, this constraint means we accept the match with the highest posterior probability, and declare the other to have no match. A similar approach can be see in the most probable maximal matching sets used by (Steorts, 2013) to match records to latent entities, and is justified within this Bayesian framework.

Record linkage is the task of identifying duplicate records across multiple data sources, often in the absence of a unique identifier (Christen, 2012). Record linkage is an increasingly important task in "data cleaning," that is used for inferential and predictive analyses in fields such as statistics, computer science, machine learning, political science, economics, precision medicine, official statistics, and others. Many statistical record linkage methods and case studies are extensions of seminal work of Fellegi and Sunter (1969) and Newcombe et al. (1959). Specifically, Fellegi and Sunter (1969) created comparison vectors for each pair of records in the data and independently classified those pairs as a match or a non-match using a likelihood ratio test (or hypothesis test). Recent work in the statistical literature has extended the aforementioned work (Winkler and Thibaudeau, 1991; Fair, 2004; Wagner et al., 2014; Gill and Goldacre, 2003).

In this paper, we consider bipartite record linkage, which is commonly known as merging two databases that contain duplication across the two databases, but do not contain any duplications within each database (Sadinle, 2017). Much of the statistical literature focuses on bipartite record linkage (Fellegi and Sunter, 1969; Jaro, 1989; Winkler, 1988; Belin and Rubin, 1995; Larsen and Rubin, 2001; Tancredi and Liseo, 2011; Herzog et al., 2007; Gutman et al., 2013; Sadinle, 2017).

The Fellegi-Sunter method and its extensions are widely popular mainly due to its simplicity and computational scalability. Despite this, like any method, it has limitations, that are well known in the literature. Specifically, the assumption of no duplications within a database implies a maximum one-to-one constraint. That is, a record from one database can be linked with only one record in the other database. Modern methods of Fellegi-Sunter either ignore this restriction (Winkler, 1988; Belin and Rubin, 1995; Larsen and Rubin, 2001), include this as a post-processing step (Jaro, 1989), or include it directly in the modeling framework (Sadinle, 2017). It is important to note that Fellegi and Sunter (1969) ignored this restriction as well. In practice, it seems that one should either directly include it in the modeling framework or correct for this in a post-processing step.

To our knowledge, Sadinle (2017) was the first to propose correcting the Fellegi and Sunter model for the one-to-one constraint. Specifically, the authors propose a Bayesian variant of the original model that satisfies two goals. First, it allows one to quantify uncertainty regarding match or non-match status. In addition, posterior distributions are readily available to help estimate point estimates of interest (and quantity such estimates). Second, under the Bayesian paradigm, the one-to-one matching constraint is easily embedded. The downside to this constraint is an added computational cost. The goal of our paper is proposing an extension of Sadinle (2017) that scales to large databases, but does not compromise much in terms of accuracy on the original case study considered by the author.

Our Contribution In this paper, we propose a Bayesian extension of Sadinle (2017) (`BRL`) through a modification to the model specification that allows for parallel computing of the linkage parameter, using hashing techniques to hasten calculations and to reduce computational complexity. Specifically, Sadinle (2017) proposes a prior distribution that strictly enforces one-to-one matching. Instead, we propose independent priors (breaking the one-to-one matching) in order to create computational speeds and gains. Second, we resolve the one-to-one matching issue via a post processing step. Third, we derive conditional distributions, a Gibbs sampler, and corresponding Bayes estimates for our proposed model. Fourth, we explore other computational speeds up based upon storage efficient indexing. Finally, we explore a sensitivity analysis and revisit the case study of the El Salvadoran conflict to understand the practical ramifications regarding our proposed modeling choices. To recognize the lineage from the original `BRL` method, we name our method *fast beta linkage* (`fabl`, pronounced "fable")

The remainder of the paper is as follows. Section **??** reviews bipartite record linkage; section **??** reviews the traditional approach of the Fellegi and Sunter framework before introducing our proposed work.

4   BIPARTITE RECORD LINKAGE AND RELATED WORK

As already mentioned, bipartite record linkage is the traditional set up studied in record linkage, where we assume there are two databases $\boldsymbol{X}_1$ and $\boldsymbol{X}_2$ such that there are duplications

across the databases but not within each databases. We assume that each database has n_1 and n_2 records, and that $n_1 \leq n_2$. Denote the number of entities simultaneously recorded by *both databases* by n_{12} such that $0 \leq n_{12} \leq n_2$. Given this set up, as explained by Steorts et al. (2016) and Sadinle (2017), this can be viewed as bipartite record linkage. This is important as it allows one to view the matching matrix or linkage structure in the following way:

$$\Delta_{ij} = \begin{cases} 1 & \text{if records } i \in \boldsymbol{X}_1 \text{ and } j \in \boldsymbol{X}_2 \text{ refer to the same entity;} \\ 0 & \text{otherwise.} \end{cases} \tag{7}$$

Bipartite matching can also be viewed as a matching labeling $\boldsymbol{Z} = (Z_1, \ldots, Z_{n_2})$ for the records in database $\boldsymbol{X}_2$ such that

$$Z_j = \begin{cases} i & \text{if records } i \in \boldsymbol{X}_1 \text{ and } j \in \boldsymbol{X}_2 \text{ refer to the same entity;} \\ n_1 + j & \text{if records } j \in \boldsymbol{X}_2 \text{ does not have a match in database } \boldsymbol{X}_1. \end{cases} \tag{8}$$

Depending on which representation is the most convenient, we can go back and forth between the two using $\Delta_{ij} = I(Z_j = i)$, where $I(\cdot)$ is the indicator function.

There are many approaches to bipartite record linkage, and we review some of the most common approaches below.

Fellegi and Sunter   As already mentioned, the most common approach is the original approach of Fellegi and Sunter (1969). Enamorado et al. (2019a) have extended this approach using computational speeds ups to allow for a much more scalable and practical approach, providing open source code for the community.

Directly Modeling the Data   Instead of relying on comparison vectors, some authors will incorporate the data directly into their bipartite record linkage model (Fortini et al., 2001; Matsakis, 2010; Liseo and Tancredi, 2011; Tancredi and Liseo, 2011; Gutman et al., 2013; Shan et al., 2020). These approaches have led to extensions regarding downstream tasks in the literature or joint models, such as regression approaches (Dalzell and Reiter, 2018; Steorts et al., 2018; Tancredi et al., 2020). There have been extensions to directly

model the data in this literature where authors consider merging multiple databases (more than two) (Steorts et al., 2016; Steorts, 2015; Zanella et al., 2016; Marchant et al., 2019; Tancredi and Liseo, 2015). We do not review this literature given that it is not directly/fairly comparable with the case of bipartite record linkage.

Classification Approaches One common way of approaching the record linkage task, especially in computer science and machine learning is using a classification approach. One uses a model (such as random forests) to classify records into matches/non-matches. When access to unique identifiers is available, one can create a training data (or reference data) set of record pairs (where the true match status is known). Finally, one can train a classifier using the reference data set to predict the match status on any record pairs that were not included in the reference data set (Bilenko et al., 2006; Christen, 2008; Ventura et al., 2015).

5 REVIEW OF THE FELLEGI SUNTER METHOD

In this section, we review notation, terminology, and the classical approach of Fellegi and Sunter.

Consider ordered record pairs $\boldsymbol{X}_1 \times \boldsymbol{X}_2$, where the size of the databases are n_1 and n_2, respectively. Denote the set of matches by $\boldsymbol{M} = \{(i,j) : i \in \boldsymbol{X}_1, j \in \boldsymbol{X}_2, \Delta_{ij} = 1\}$. Denote the set of non-matches $\boldsymbol{U} = \{(i,j) : i \in \boldsymbol{X}_1, j \in \boldsymbol{X}_2, \Delta_{ij} = 0\}$. An alternative way of merging two databases and removing duplications can be viewed as identifying the sets of matches/non-matches, or rather $\boldsymbol{M}, \boldsymbol{U}$. In the literature, it is well known that record pairs that are estimated as matches are known as links; record pairs that are estimated as non-matches are known as non-links. Fellegi and Sunter (1969) proposed using all-to-all pairwise comparison of records to estimate the matches/non-matches.

5.1 Comparison Data

In this section, we review the notion of *comparison data*. Intuitively, records that are co-referent or rather refer to the same entity should be similar; records that are non co-referent and refer to different entities should not be similar. A common way to encode this is using a

comparison vector $\boldsymbol{\gamma}_{ij}$ which is computed for each record pair (i, j) in $\boldsymbol{X}_1 \times \boldsymbol{X}_2$. Denote F as the criteria for compare the records such that $\boldsymbol{\gamma}_{ij} = (\gamma_{ij}^1, \gamma_{ij}^2, \ldots, \gamma_{ij}^f, \ldots \gamma_{ij}^F)$. In most cases, F represents to a single comparison per feature that the databases share, such as gender. The simplest way to compare two records is check for agreement or disagreement, and this is commonly used for nominal variables such as gender. For more complex measurements, such as textual or numeric feature information, we can take into account partial agreement patterns using common string metrics, such as the Edit, Jaro, or Jaro-Winkler distance functions. This allows us to divide our records into a set of similarity values into different levels of disagreement (Bilenko and Mooney, 2006; Elmagarmid et al., 2007).

Let $\mathcal{S}_f(i, j)$ denote a general similarity measure for feature f of records i and j, where the range of $\mathcal{S}_f$ can be divided into $L_f + 1$ intervals denoted by $I_{f0}, I_{f1}, \ldots, I_{fL_f}$, representing disagreement levels. Following convention, I_{f0} represents the highest level of agreement (inclusive of complete agreement) and I_{fL_f} represents the highest level of disagreement (including complete disagreement). Thus, we can construct comparison vectors in the following way:

$$\gamma_{ij} = \ell \text{ if } \mathcal{S}_f(i, j) \in I_{f\ell}.$$

The choice of $I_{f\ell}$ are application specific, which we discuss in our simulation and case study.

5.2 Blocking

In practice, it is not feasible to make all-to-all record comparisons as the computationally complexity is of the order $O(n_1 \times n_2)$. The most common solution is utilize blocking, which places similar records into partitions, clusters, or "blocks" to reduces this computational burden. Blocking can be deterministic, such as reducing the dimensionality of the comparison space on a highly reliable feature, such a gender. Blocking can be probabilistic, such as utilizing machine learning methods, known as hashing, which seek to use probabilistic functions (with low computational cost) to place similar records into the same bin. We propose hashing base methods for Bayesian bipartite record linkage, illustrating their ability to improve upon computational scalability. We emphasize that depending on the size of the problem, our proposed methodology can be used with or without hashing.

5.3 Fellegi-Sunter Model

In this section, we review the Fellegi-Sunter model. While we have introduced comparison data, $\boldsymbol{\gamma}_{ij}$, this is not sufficient to determine whether a record is a match/non-match given errors that naturally occur in the data. This motivated Fellegi and Sunter (1969) to consider the following likelihood ratio

$$w_{ij} = \log \frac{P(\gamma_{ij} \mid \Delta_{ij} = 1)}{P(\gamma_{ij} \mid \Delta_{ij} = 0)} \tag{9}$$

as a weight to estimate if record pairs are a match/non-match. Assume $\boldsymbol{\Gamma}_{ij}$ is a random vector, whose distribution depends on Δ_{ij}. Then equation 9 is a representative value of $\boldsymbol{\Gamma}_{ij} \mid \Delta_{ij}$. The ratio will be large if we favor the pair being a match. This is a problem in practice as transitive closures are violated. For example, suppose records i and j in $\boldsymbol{X}_1$ are extremely similar but are known to be non-matches. Suppose that both records i and j in $\boldsymbol{X}_1$ are very similar to record k in $\boldsymbol{X}_2$. By transitive closures, this implies that i and j are a match, which cannot be true by assumption.

Jaro (1989) proposed addressing this issue via an optimization approach as follows:

$$\max_{\Delta} \sum_{i=1}^{n_1} \sum_{j=1}^{n_2} w_{ij} \Delta_{ij} \quad \text{subject to} \quad \Delta_{ij} \in \{0, 1\}; \tag{10}$$

$$\sum_{i=1}^{n_1} \Delta_{ij} \leq 1, j = 1, 2, \ldots n_2; \text{ and}$$

$$\sum_{i=1}^{n_2} \Delta_{ij} \leq 1, j = 1, 2, \ldots n_1.$$

Equation 10 ensures that each record of $\boldsymbol{X}_2$ only matches with at most one record in $\boldsymbol{X}_1$ (and the reverse condition). This is commonly known as the *maximum-weight bipartite matching problem* or a *linear sum assignment problem*. The output of this algorithm is a bipartite matching that maximizes the sum of the weights among matched pairs, where the pairs are not matched are assumed to be non-matches. Jaro (1989) provided no theoretical justification for using this approach, however, Sadinle (2017) recently showed that under certain conditions, equation 10 is the maximum likelihood estimate.

5.4 Model Estimation

In our review, we have assumed one knows $P(\cdot \mid \Delta_{ij} = 1)$ and $P(\cdot \mid \Delta_{ij} = 0)$, however, these are typically not known in practice and must be estimated. Assume the comparison vectors are independent given the bipartite matching and assume that the match status of the record pairs are independent. It then follows that we can model the comparison data using mixture models in the following way:

$$\Gamma_{ij} = \gamma_{ij} \mid \Delta_{ij} = 1 \overset{iid}{\sim} \mathcal{M}(\boldsymbol{m}), \tag{11}$$
$$\Gamma_{ij} = \gamma_{ij} \mid \Delta_{ij} = 0 \overset{iid}{\sim} \mathcal{U}(\boldsymbol{u}),$$
$$\Delta_{ij} \overset{iid}{\sim} \text{Bernoulli}(p),$$

where p is the proportion of records that match. Note that $\boldsymbol{m}$ and $\boldsymbol{u}$ are vectors of the match/non-match parameters. One common way to estimate the unknown parameters of the aforementioned model is using the EM algorithm.

One criticism of using mixture models as proposed by Jaro (1989) is the fact that the decision rule will lead to "many-to-many assignments (breaking transitive closures). One way of resolving this issue to incorporate the constraint directly into the model (Fortini et al., 2001; Matsakis, 2010; Liseo and Tancredi, 2011; Tancredi and Liseo, 2011; Larsen, 2005b, 2012; Gutman et al., 2013; Sadinle, 2017), which is appealing as one then obtains uncertainty quantification under the Bayesian paradigm, but often has to sacrifice scaling to very large databases. On the other hand, one can turn to post-processing methods in practical situations. In this paper, we attempt to bridge the best of both approaches. Specifically, we propose a model in the spirit of Sadinle (2017) and Enamorado et al. (2019b). We relax the one-to-one matching constraint proposed by Sadinle (2017) for purely computational reasons, and resolve this in a post-processing step. In addition, we propose our own computational speeds up and those in the spirit of Enamorado et al. (2019b). Our framework is then embarassingly parallel, available in the `fabl` package, and evaluated on a real case study and simulation studies illustrating that any loss in accuracy due to removing the constraint is minimal.

6 PROPOSED METHODOLOGY

In this section, we propose the fast beta linkage method (fabl) for bipartite record linkage, extending the work of Sadinle (2017). Specifically, we review the general approach of Sadinle (2017) and then provide our extensions for scalability.

6.1 Beta Linkage Model

In this section, we review the Beta Linkage Model, which can be written as follows:

$$\boldsymbol{\Gamma}_{ij} \mid Z_j = i \overset{iid}{\sim} \mathcal{M}(\boldsymbol{m}) \tag{12}$$

$$\boldsymbol{\Gamma}_{ij} \mid Z_j \neq i \overset{iid}{\sim} \mathcal{U}(\boldsymbol{u})$$

$$\boldsymbol{Z} \sim \mathcal{B},$$

where $\mathcal{B}$ represents a prior on the space of bipartite matchings and $\mathcal{M}(\boldsymbol{m}), \mathcal{M}(\boldsymbol{u})$ are models for the comparison vectors of matches and non-matches.

6.2 New Model Specification

In this section, we provide the specification for the models of $\mathcal{M}(\boldsymbol{m})$ and $\mathcal{U}(\boldsymbol{u})$, which follows from assuming that the features are conditionally independent of the level of agreement of another. The likelihood can be written as follows:

$$\mathcal{L}(\boldsymbol{Z}, \Phi \mid \boldsymbol{\Gamma}) = \prod_{j=1}^{n_1} \prod_{i=1}^{n_2} \left[\prod_{f=1}^{F} \prod_{l=1}^{L_f} m_{fl}^{I(Z_j=i)} u_{fl}^{I(Z_j \neq i)} \right]^{I(\gamma_{ij}^f = l)}, \tag{13}$$

where $m_{f\ell} = P(\Gamma_{ij}^f = \ell \mid Z_j = i)$ denotes the probability of a match have level ℓ of disagreement in feature f and $u_{f\ell} = P(\Gamma_{ij}^f = \ell \mid Z_j \neq i)$ represents the same probability for non-matches. In addition, denote $\boldsymbol{m}_f = (m_{f1}, \ldots, m_{fL_f})$, $\boldsymbol{u}_f = (u_{f1}, \ldots, u_{fL_f})$, $\boldsymbol{m} = (m_1, \ldots m_F)$, $\boldsymbol{u} = (u_1, \ldots u_F)$, and $\Phi = (\boldsymbol{m}, \boldsymbol{u})$.

Next, we allow for missing values or rather missing comparisons for corresponding record pairs. As commonly done in the record linkage literature, we assume that these occur missing at random [CITE: Little and Rubin (2002), Sadinle (2017), Steorts et al (2016)].

The likelihood becomes

$$\mathcal{L}(\boldsymbol{Z}, \Phi \mid \boldsymbol{\Gamma}^{obs}) = \prod_{j=1}^{n_1} \prod_{i=1}^{n_2} \left[\prod_{f=1}^{F} \prod_{l=1}^{L_f} m_{fl}^{a_{f\ell}(\boldsymbol{Z})} u_{fl}^{a_{f\ell}(\boldsymbol{Z})} \right], \tag{14}$$

where

$$a_{f\ell}(\boldsymbol{Z}) = \sum_{i,j} I_{obs}(\gamma_{ij}^f) I(\gamma_{ij}^f = \ell) I(Z_j = i) \tag{15}$$

$$b_{f\ell}(\boldsymbol{Z}) = \sum_{i,j} I_{obs}(\gamma_{ij}^f) I(\gamma_{ij}^f = \ell) I(Z_j \neq i), \tag{16}$$

where I_{obs} denotes if a value is observed or not. Next, for a particular matching configuration $\boldsymbol{Z}$, $a_{f\ell}(\boldsymbol{Z})$ and $b_{f\ell}(\boldsymbol{Z})$ denote the number of matches and non-matches with observed disagreement level ℓ for comparison feature f. Using this parameterization, we specify independent conjugate priors for all features f:

$$\boldsymbol{m}_f \sim \text{Dirichlet}(\alpha_{f0}, \dots, \alpha_{fL_f}) \tag{17}$$

$$\boldsymbol{u}_f \sim \text{Dirichlet}(\beta_{f0}, \dots, \beta_{fL_f}). \tag{18}$$

6.3 Fast Beta Prior

In this section, we deviate from the approach of Sadinle (2017), and propose a more scalable prior on the matching labelings $\boldsymbol{Z} = (Z_1, Z_2, \dots Z_{n_2})$, where recall $Z_j \in \{1, 2, 3, n_1, n_1 + j\}$. While the rest of the model has been proposed previously, to our knowledge, this particular prior has not been proposed before.

$$Z_j | \lambda = \begin{cases} \frac{1}{n_1} \lambda & z_j \leq n_1; \\ 1 - \lambda & z_j = n_1 + j \end{cases}$$

$$\lambda \sim \text{Beta}(\alpha_\lambda, \beta_\lambda)$$

Intuitively, this set of priors says that record $j \in \boldsymbol{X}_2$ has some match in $\boldsymbol{X}_1$ with probability λ, and that each record $i \in \boldsymbol{X}_1$ is equally likely to be that match. λ itself is given a prior distribution to reflect prior beliefs about the overall rate of matchingness in

the data. One choice of uninformative prior might be $\lambda \sim \text{Beta}(1,1)$, which corresponds to a prior belief that nonmatches and matches are equally likely, and another might be $\lambda \sim \text{Beta}\left(1, \frac{1}{n_1}\right)$, which corresponds to a uniform prior on the labeling of $\mathbf{Z}$.

There is a clear relationship between our proposed model and that of Sadinle (2017). Sadinle (2017) constructs a prior distribution on the entire $\mathbf{Z}$ vector, called the *beta distribution for bipartite matching*, given by

$$P(\mathbf{Z}|\alpha_\pi, \beta_\pi) = \frac{(n_1 - n_{12}(\mathbf{Z}))!}{n_1!} \frac{\text{B}(n_{12}(\mathbf{Z}) + \alpha_\pi, n_2 - \mathbf{Z}) - \beta_\pi])}{\text{B}(\alpha_\pi, \beta_\pi)}$$

where $\text{B}(\cdot, \cdot)$ represents the Beta function. This prior induces a Gibbs sampler that strictly enforces one-to-one matching, previously matched records from the set of candidate records when sampling Z_j. This creates a dependency that makes the sampler *inherently serial*.

On the other hand, we propose independent priors for each Z_j, creating a sampler that is *embarassingly parallel*, allowing for significant computational gains. More importantly, since only the agreement pattern of Z_j is used for calculations within the Gibbs sampler, and not the particular record label, this sampling is only needed at the level of the unique agreement patterns, which offers additional computational savings. In doing so however, we thereby weaken the one-to-one requirement from BRL; our sampler does ensure that each record in $\boldsymbol{X}_2$ can be matched to at most one record in $\boldsymbol{X}_1$, but allows for the possibility that multiple records in $\boldsymbol{X}_2$ match to the same record in $\boldsymbol{X}_1$.

6.4 Gibbs Sampler

We work with the following factorization of the joint posterior distribution:

Brian: The Gibbs sampler has some minor typos. It should be written in terms of $\mathbf{Z}, \Phi, \lambda \mid \gamma_{obs}$. Also, just make the rest of the notation match please:)

$$p(\mathbf{Z}, \mathbf{m}, \mathbf{u}, \lambda | \Gamma) \propto \mathcal{L}(\mathbf{Z}, \Phi \mid \boldsymbol{\gamma}^{obs}) p(\mathbf{Z}|\lambda) p(\Phi) p(\lambda)$$

$$\propto \prod_{j=1}^{n_2} \prod_{i=1}^{n_1} \left[\prod_{f=1}^{F} \prod_{l=1}^{L_f} m_{fl}^{I(Z_j=i)} u_{fl}^{I(Z_j \neq i)} \right]^{I(\gamma_{ij}^f=l)}$$

$$\times \prod_{f=1}^{F} \prod_{l=1}^{L_f} m_{fl}^{\alpha_{fl}-1} \times \prod_{f=1}^{F} \prod_{l=1}^{L_f} u_{fl}^{\beta_{fl}-1}$$

$$\times \prod_{j=1}^{n_B} \left[I(Z_j \leq n_1) \frac{1}{n_1} \lambda + I(Z_j > n_1)(1 - \lambda) \right]$$

This factorization leads to following Gibbs Sampler:

<u>Sample $\mathbf{m}^{(s+1)}$ $\mathbf{u}^{(s+1)}|\Gamma, \mathbf{Z}^{(s)}$</u>: The $\mathbf{m}$ and $\mathbf{u}$ parameters are updated through standard multinomial-dirichlet mechanics. Thus we have

$$\mathbf{m}_f|\mathbf{Z}, \Gamma \sim \text{Dirichlet}(\alpha_{f1}(\mathbf{Z}), \ldots, \alpha_{fL_f}(\mathbf{Z}))$$

$$\mathbf{u}_f|\mathbf{Z}, \Gamma \sim \text{Dirichlet}(\beta_{f1}(\mathbf{Z}), \ldots, \beta_{fL_f}(\mathbf{Z}))$$

where $\alpha_{fl}(\mathbf{Z}) = \alpha_{fl} + \sum_{i,j} I(\gamma_{ij}^f = l) I(z_j = i)$ and $\beta_{fl}(\mathbf{Z}) = \beta_{fl} + I(\gamma_{ij}^f = l) I(z_j \neq i)$.

<u>Sample $\lambda^{(s+1)}|\mathbf{Z}^{(s)}$</u>: As a function of λ, the linkage structure parameter $\mathbf{Z}$ is sequence of successes (when $z_j < n_A + 1$) and failures (when $z_j = n_A + 1$), and therefore $p(\mathbf{Z}|\lambda) = \mathcal{L}(\lambda|\mathbf{Z})$ is determined only by the number of duplicates $n_{12}(\mathbf{Z}) = \sum_{i=1}^{n_2} I(z_j < n_1 + 1)$ encoded by $\mathbf{Z}$. Thus we have

$$p(\lambda|\mathbf{Z}) \propto p(\mathbf{Z}|\lambda) p(\lambda)$$

$$\propto \lambda^{n_{12}(\mathbf{Z})} (1 - \lambda)^{n_2 - (n_{12}(\mathbf{Z}))} \lambda^{\alpha_\lambda - 1} (1 - \lambda)^{\beta_\lambda - 1}$$

$$\propto \lambda^{n_{12}(\mathbf{Z}) + \alpha_\lambda - 1} (1 - \lambda)^{n_1 - n_{12}(\mathbf{Z}) + \beta_\lambda - 1}$$

$$\implies \lambda^{(s+1)}|\mathbf{Z}^{(s+1)} \sim \text{Beta}(D + \alpha_\lambda, n_2 - n_{12}(\mathbf{Z}) + \beta_\lambda)$$

Sample $\mathbf{Z}^{(s+1)}|\Gamma, \mathbf{m}^{(s+1)}, \mathbf{u}^{(s+1)}, \lambda^{(s+1)}$: Because we sample Z_j independently of all other $Z_{j'}$, we use only the full conditional for an individual Z_j. Let $\Gamma_{\cdot j}$ denote the set of n_A comparison vectors with $j \in B$, and note that as a function of Z_j, the likelihood $p(\Gamma_{\cdot j}|Z_j, \mathbf{m}, \mathbf{u}) = \mathcal{L}(Z_j|\Gamma_{\cdot j}, \mathbf{m}, \mathbf{u})$ is a discrete distribution with probabilities proportional to

$$p(\Gamma_{\cdot j}|Z_j = z_j, \mathbf{m}, \mathbf{u}) \propto \prod_{i=1}^{n_1} \left[\prod_{f=1}^{F} \prod_{l=1}^{L_f} m_{fl}^{I(Z_j=i)} u_{fl}^{I(Z_j \neq i)} \right]^{I(\gamma_{ij}^f = l)}$$

$$\propto \prod_{i=1}^{n_1} \left(\prod_{f=1}^{F} \prod_{l=1}^{L_f} \frac{m_{fl}}{u_{fl}} \right)^{I(z_j=i, \gamma_{ij}^f = l)}$$

$$= \begin{cases} w_{ij} & z_j \leq n_1; \\ 1 & z_j = n_1 + j \end{cases}$$

where $w_{ij} = \left(\frac{\prod_{f=1}^{F} \prod_{l=1}^{L_f} m_{fl}}{\prod_{f=1}^{F} \prod_{l=1}^{L_f} u_{fl}} \right)^{I(\gamma_{ij}^f = l)} = \frac{P(\boldsymbol{\gamma_{ij}}|Z_j=i)}{P(\boldsymbol{\gamma_{ij}}|Z_j \neq i)}$. The interested reader should note that these are precisely the likelihood ratios used in the Fellegi-Sunter model to classify matches and non-matches, and we therefore refer to w_{ij} as the *Fellegi Sunter weights*.

With the likelihood in this form, we can derive the full conditional

$$p(Z_j|\Gamma_{\cdot j}, \mathbf{m}, \mathbf{u}, \lambda) \propto p(\Gamma_{\cdot j}|Z_j, \mathbf{m}, \mathbf{u})P(Z_j|\lambda)$$

$$\propto \left(\sum_{i=1}^{n_A} w_{ij} \mathbf{1}_{z_j=i} + \mathbf{1}_{z_j=n_A+1} \right) \left(\lambda \sum_{i=1}^{n_A} \frac{1}{n_A} \mathbf{1}_{z_j=i} + (1-\lambda)\mathbf{1}_{z_j=n_A+1} \right)$$

$$= \frac{\lambda}{n_A} \sum_{i=1}^{n_A} w_{ij} \mathbf{1}_{z_j=i} + (1-\lambda)\mathbf{1}_{z_j=n_A+1}$$

$$\implies Z_j^{(s+1)}|\mathbf{m}, \mathbf{u}, \Gamma, \lambda \propto \begin{cases} \frac{\lambda}{n_A} w_{ij} & z_j \leq n_A; \\ 1 - \lambda & z_j = n_A + 1 \end{cases}$$

In order to make fair comparisons against the (Sadinle, 2017) model, we integrate over the posterior of λ and rearrange terms to produce the final full conditional:

$$p\left(Z_j^{(s+1)} = i|\mathbf{m}, \mathbf{u}, \mathbf{Z}^{(\mathbf{s})}\right) \propto \begin{cases} w_{ij} & i \leq n_A; \\ n_A \frac{n_B - D + \beta_\lambda}{D + \alpha_\lambda} & i = n_A + 1 \end{cases}$$

6.5   Bayes Estimate

We calculate a Bayes estimate $\hat{\mathbf{Z}}$ for the linkage parameter $\mathbf{Z}$ by assigning different positive losses to different types of errors, and minimizing posterior expected loss. We adopt the loss function proposed in Sadinle 2017, in which $\hat{Z}_j \in \{1, \ldots, n_A + 1, R\}$, with R representing the option to leave the matching undetermined by the model. Specifically, we have

$$L(\hat{Z}_j, Z_j) = \begin{cases} 0 & \text{if } Z_j = \hat{Z}_j; \\ \theta_R, & \text{if } \hat{Z}_j = R; \\ \theta_{10}, & \text{if } Z_j \leq n_A, \hat{Z}_j = n_A + 1; \\ \theta_{01}, & \text{if } Z_j = n_A + 1, \hat{Z}_j \leq n_A; \\ \theta_{11}, & \text{if } Z_j, \hat{Z}_j \leq n_A, Z_j \neq \hat{Z}_j; \end{cases}$$

Here, θ_R is the loss from not making a decision on the linkage status, θ_{10} is the loss from a false non-match, θ_{01} is the loss from a false match, and θ_{11} is the loss from the special case of a false match in which the record has a true match other than the one estimated by the model.

We seek to minimize the posterior expected loss, given by $\mathbb{E}[L(\hat{\mathbf{Z}}, \mathbf{Z})|\Gamma^{obs}] = \sum_{\mathbf{Z}} L(\hat{Z}_j, Z_j)P(\mathbf{Z}|\Gamma^{obs})$. In this paper, we adopt losses $\theta_R = \infty, \theta_{10} = 1, \theta_{01} = 1, \theta_{11} = 2$, inducing the intuitive decision rule

$$\hat{Z}_j = \begin{cases} i, & \text{if } P(Z_j = i|\Gamma) > \frac{1}{2}; \\ 0, & \text{otherwise}; \end{cases}$$

For a more in-depth explanation of this function and the induced Bayes estimate, see (Sadinle, 2017).

27

Since our Gibbs procedure does not strictly enforce one-to-one matching, it is possible for this Bayes estimate to link multiple records in $\boldsymbol{X_2}$ to one record in $\boldsymbol{X_1}$. It is often desirable to see such matches, because they provide evidence of model misspecification; specifically, large numbers of "many to one" matches indicate that the assumption of no duplicates within files may not be reasonable. We will explore this issue more deeply through the El Salvador homicide case study.

To achieve a Bayes estimate that fulfills one-to-one matching requirement, we simply minimize the expected loss subject to the constraint that $\hat{Z}_j \neq \hat{Z}_{j'}$ for all $j \neq j'$. In the event that two matches. In the event that we have two records j and j' such that both $P(Z_j = i|\Gamma) > \frac{1}{2}$ and $P(Z_{j'} = i|\Gamma) > \frac{1}{2}$, this constraint means we accept the match with the highest posterior probability, and declare the other to have no match. A similar approach can be see in the most probable maximal matching sets used by (Steorts, 2013) to match records to latent entities, and is justified within this Bayesian framework.

7   EFFICIENT AND SCALABLE IMPLEMENTATION

In this section, we propose efficient and scalable proposals using hashing based methods that are similar in spirit to (**?**), a fast and scalable implementation of the Fellegi-Sunter model. Specifically, one way we can improve our computational efficiency is by recognizing that record pairs contribute to posterior calculations only through the agreement pattern of the γ_{ij} vector. To make this more precise, let $\mathcal{H}$ be the set of unique agreement patterns in the data, let $P = |\mathcal{H}|$ denote the total number of unique agreement patterns. Observe that P is bounded above by $\prod_{f=1}^{F} L_f$, and that this bound does not depend on n_1 or n_2. Prior to processing the data, we identify all P patterns in $\mathcal{H}$ and enumerate them as follows: $h_1, \ldots, h_P$, which are called *hashed values.* Next, we map record pairs to their corresponding hashed value. More specifically, when the record pair (i, j) exhibits the p^{th} agreement pattern, we say $(i, j) \in h_p$. Whenever possible, we conduct calculations over these P agreement patterns, instead of the typical $n_1 \times n_2$ record pairs.

7.1 Data Representation, Hashing, and Storage

In this section, we describe the hashing function we propose.

First, we hash record pairs of the same agreement pattern to unique integer values. (**?**) accomplished this efficiently through the hashing function

$$\tilde{\gamma}_{ij} = \sum_{f=1}^{F} I(\gamma_{ij}^f > 0) 2^{\gamma_{ij}^f + I(k>1) \sum_{e=1}^{k-1}(L_e - 1)}$$

This function maps each agreement pattern to a unique integer, allowing us to store a scalar quantity instead of an entire vector for each record pair. For computational ease, we then map these integers to sequential integers from $1, \ldots, P$.

The classic Fellegi Sunter method represents the γ_{ij} comparison as a vector of length F, with each component γ_{ij}^f taking on values in $\{1, \ldots, L_f\}$. To ease our computational burden, we instead use a *one hot encoding* of the comparison vector. For example, if $L_1 = L_2 = 2$ and $L_3 = 3$, then $\gamma_{ij} = (2, 1, 3)$ under the classical framework becomes $\gamma_{ij} = (0, 1, 1, 0, 0, 0, 1)$ under our framework. This is a bijective transformation that does not change the meaning of the data, but this representation eases calculations and posterior updates. This is also the form the data takes in the BRL package in R.

In the classic Fellegi Sunter framework, $\boldsymbol{\Gamma}$ is a $n_1 n_2 \times F$ matrix, with each row providing the comparison vector for a different (i, j) pair. In contrast, we do not store these comparison vectors themselves, but instead only the hashed value h_p corresponding to the agreement pattern of the (i, j) pair. We store this information in a nested list $\tilde{\boldsymbol{\Gamma}}$ where the p^{th} component of the j^{th} list contains a vector of records in $\boldsymbol{X}_1$ that share agreement pattern p with record $j \in \boldsymbol{X}_2$. For each p, we also calculate $H_p = \sum_{i=1}^{n_1} \sum_{j=1}^{n_2} \mathbf{1}_{(i,j) \in h_p}$, the total instances of agreement pattern p throughout the data, and also for each j, we calculate $H_{p_j} = \sum_{i=1}^{n_1} \mathbf{1}_{(i,j) \in h_p}$ the instances of agreement pattern p among the comparison vectors between record $j \in \boldsymbol{X}_2$ and each of the n_1 records in $\boldsymbol{X}_1$.

The hashing procedure described above considerably reduces the memory needed to store the comparison information. That is, instead of storing $n_1 \times n_1$ comparison vectors, which are relatively long under either the Fellegi Sunter or our modified framework, we only store the P unique vectors, and then $n_1 \times n_2$ scalar quantities relating record pairs to those

vectors. However, even storing these $n_1 \times n_2$ scalar labels can become burdensome with large data. Worse, the overwhelming majority of these labels relate to record pairs that are clear non-matches.

To address this issue, we propose a new method called *storage efficient indexing* (SEI). In standard indexing, one decides a certain criteria that they expect all true matching pairs to satisfy. In addition, one decides a priori to label any record pairs that do not meet that criteria as non-matches. For example, one might only consider pairs with a certain similarity score on a field deemed to be important (like first name), or only pairs with exact matching on a specified number of fields. While generally chosen to be be quite loose, establishing these criteria requires knowledge of the problem and invites room for human error. To improve upon this, we propose to reduce the comparison space and reduce the storage requirements while avoiding the drawbacks of standard indexing.

Observe that all records of the same agreement pattern have the same probability when sampling Z_j. Therefore we know that records belonging to an h_p such that H_{p_j} is large are very unlikely to be sampled consistently enough to be deemed a match through the Bayes estimate, even without considering the form of the agreement pattern itself.

In SEI, rather than store all of these unlikely record labels, we choose to store only a small number R of them. Posterior calculations still attribute the appropriate weight to all records through the summary statistics H_p, and H_{p_j}. Rather than storing $n_1 \times n_2$ record labels, SEI allows us to store at most $n_2 \times P \times R$ labels, regardless of how large n_1 might be.

Lastly, for large data, we can partition the two datasets $\boldsymbol{X}_1$ and $\boldsymbol{X}_2$ into smaller blocks $\{\boldsymbol{X}_{1m}\}$ and $\{\boldsymbol{X}_{2m}\}$ for more manageable computations. On a single machine, we can read-in data sequentially, conduct hashing, compress information through SEI, and delete the original data from memory before continuing with the next chunk of data. With multiple cores or multiple machines, this can be done in parallel. Thus, the combination of hashing, SEI, and partitioning allows us to conduct linkage tasks over much larger datasets.

7.2 Efficient Posterior Inference

Updating $\mathbf{m}$ and $\mathbf{u}$: After receiving matching statuses from $\mathbf{Z}$, the Sadinle method calculates $\alpha_{fl}(\mathbf{Z})$ and $\beta_{fl}(\mathbf{Z})$ for each field and level. This constitutes $2 \times \sum L_f$ many summations over $n_A \times n_B$ quantities, and becomes computationally burdensome with large data. In contrast, we recognize that each unique agreement pattern contributes to the posterior $\alpha(\mathbf{Z})$ and $\beta(\mathbf{Z})$ vectors in the same way. In fact, if we denote $H_p^m = \sum_{j=1}^{n_B} \mathbf{1}_{(Z_j,j) \in h_p}$ to be the number of matching record pairs with agreement pattern p, then the contribution of pairs of pattern p to the $\alpha(\mathbf{Z})$ vector is simply $H_p^m \times h_p$. Thus our posterior update for the α vector is simply $\alpha(\mathbf{Z}) = \alpha_0 + \sum_{p=1}^{P} H_p^m \times h_p$. Then, we can easily calculate H_p^u, the number of nonmatching record pairs of agreement pattern p, by subtracting the number of matching pairs from the total present in the data; that is $H_p^u = H_p - H_p^m$. From this, we can update our β parameter through $\beta(\mathbf{Z}) = \beta_0 + \sum_{p=1}^{P} H_p^u \times h_p$. Note that these constitute P many summations over n_B quantities, and thus avoid the $n_A \times n_B$ summation from the original method.

Updating $\mathbf{Z}$: Although sampling Z_j from a the full conditional provided earlier is conceptually straightforward, it becomes computational burdensome when n_A is larger. The reader can confirm that sampling a value from a large set of unequal probabilities becomes difficult in most programming languages. To speed up computation, we break this sampling step into two simpler steps. First, we calculate the Fellegi Sunter weight w_p associated with each unique pattern and sample the agreement pattern between j and its potential match. Second, we sample the record label uniformly among records associated with that agreement pattern for that particular $j \in B$. More concretely, define $h(Z_j)$ to be the agreement pattern between j and its potential match, and say $h(Z_j) = h_{P+1}$ when $Z_j = n_A + 1$. Then,

$$P\left(h\left(Z_j^{(s+1)}\right) = p | \mathbf{m}, \mathbf{u}, \mathbf{Z}^{(s)}\right) \propto \begin{cases} w_p \times H_{p_j} & p \leq P; \\ n_A \frac{n_B - D + \beta_\lambda}{D + \alpha_\lambda} & p = P+1 \end{cases}$$

$$P\left(Z_j^{(s+1)} = i \middle| h\left(Z_j^{(s+1)}\right) = p\right) = \begin{cases} \frac{1}{H_{p_j}} & (i,j) \in h_p \\ 0 & \text{otherwise} \end{cases}$$

Lastly, we recognize that all posterior updates are governed by the agreement patterns of the record pairs rather than the record labels themselves. Thus we complete the entire Gibbs procedure first at the level of the P agreement patterns with the first equation above. After, we can back-fill the records corresponding to the agreement patterns through the second equation. Sampling uniformly is computationally simple even for large sets of candidate records, but this step can also be parallelized when working with large data.

To aid the reader, we provide summary of the `fabl` method through pseudocode:

Algorithm 1 Summary of fabl algorithm

1: **procedure** HASHING AND PREPROCESSING
2:      Partition files A and B into chunks $\{A_I\}, \{B_J\}$
3:      **for** each I, J **do**
4:          Create comparison vectors between A_I and B_J
5:          Hash results and calculate summary statistics
6:          Use SEI to reduce memory usage
7:      **end for**
8:      Synthesize results across pairings
9: **end procedure**
10: **procedure** GIBBS SAMPLING
11:     Initialize m, u, and Z parameters
12:     **for** $t \in \{1, \ldots, T\}$ **do**
13:         Sample $m^{t+1}|Z^t, \Gamma$ and $u^{t+1}|Z^t, \Gamma$
14:         Sample $H\left(Z^{t+1}\right)|m^{t+1}, u^{t+1}, \Gamma$ ▷ Sample agreement pattern, not record
15:     **end for**
16:     Sample $Z|H(Z), \Gamma$ ▷ Fills in record label based on agreement pattern
17: **end procedure**

7.3 Computational Complexity

The computational complexity of `fabl` is determined by its two steps: constructing comparison vectors and Gibbs sampling. The computational complexity of all pairwise comparisons

across files $\mathbf{X}_1$ and $\mathbf{X}_2$ is $O(Fn_1n_2)$. TALK TO BEKA. I CAN DO THE SAME TRICKS AS TED HERE TO REDUCE COMPLEXITY. Since `fabl` is able to compute these comparisons in parallel, we can split these computations across B equally sized partitions, so the complexity becomes $O(\frac{F}{B}n_1n_2)$.

Without hashing, the computational complexity of updating the $\mathbf{m}$ and $\mathbf{u}$ parameters is $O(Fn_1n_2)$. However, by doing calculations over the agreement patterns rather than the individual record, hashing reduces the complexity to $O(FP)$.

The complexity of updating Z sequentially at the record level is $O(n_1n_2)$. With hashing, we split this sampling into two steps; first we sample that agreement pattern of the match with complexity $O(n_2P)$, and then we sample the record that exhibits that pattern with complexity $O(n_2)$. Thus we have reduced the complexity of the Gibbs sampler under `BRL` from $O(Fn_1n_2)$ to $O(n_2P)$.

Overall, the computational complexity under the complete `fabl` framework is $O(\frac{F}{B}n_1n_2)+O(n_2P)$.

8   SIMULATION STUDIES

In this section, we replicate a simulation study from **?** to compare `fabl` against `BRL` on several simulated datasets with varying amounts of error and overlap. We use first name, last name, age, and occupation for this linkage, and create comparison vectors according to the default settings of the `compareRecords` function from the `BRL` package. See Table 1 for details on the construction of comparison vectors. Each simulation identifies duplicated individuals between two datasets, each with 500 records. We conduct linkage when matching records exhibit 1, 2, and 3 errors across the four fields, and when there are 50, 250, and 450 individuals in common across datasets. We use flat priors for all m and u parameters, run the Gibbs Sampler for 1000 iterations, and discard the first 100 as burn-in.

8.1   Precision, Recall, and F-measure

In this section, we compare `fabl` to `BRL` in terms of precision, recall, and f-measure, which are commonly defined record linkage evaluation metrics that are defined in (Christen, 2012).

		Level of Disagreement			
Fields	Similarity	1	2	3	4
First and Last Name	Levenstein	0	(0, .25]	(.25, .5]	(5, .1]
Age and Occupation	Binary	Agree	Disagree		

Table 1: Construction of comparison vectors for accuracy study with simulated datasets

In cases when there are only one or two errors in matching records, and in cases with low to moderate duplication across records, we see that `fabl` provides approximately equivalent accuracy as `BRL`. We find that our method only has weakened performance in the most extreme scenario of very high errors and very high overlap across files. In these situations, `BRL` is removing large numbers of records from consideration throughout the Gibbs Sampler, making its implementation most different from `fabl`. However, such extreme linkage tasks, which such high amounts of errors and overlap, are extremely rare in practice.

8.2   Accuracy under Partial Estimates

Lastly, we review the performance of `fabl` and `BRL` using loss functions that allow partial estimates of the linkage parameter. By leaving $\theta_{10} = \theta_{01} = 1$ and $\theta_{11} = 2$, but setting $\theta_R = 0.1$, we allow the model to decline to decide a match for certain records, with nonassignment being 10% as costly as a false match. In this context, we are no longer focused on finding all true matches, but rather protecting against false matches. Thus, instead of recall, we use the *negative predictive value* (NPV), defined as the proportion of non-links that are actual non-matches. Mathematically, $NPV = \sum_{j=1}^{n_2} I(\hat{Z}_j = Z_j = n_1 + j)/\sum_{j=1}^{n_2} I(\hat{Z}_j = n_1 + j)$. We continue to use the precision, which is renamed the *positive predictive value* in this context. Lastly, we also report that rejection rate (RR), or how often the model declines to make a linkage decision, defined as RR $= \sum_{j=1}^{n_2} I(\hat{Z}_j = R)$.

We see that `fabl` maintains equivalently strong PPV as `BRL` across all linkage settings. However, with high amounts of overlap and high amounts of error, the rejection rate under `fabl` rises, leading to a decrease in NPV. This is reasonable because large numbers of errors lead to situations where one record in $\mathbf{X_2}$ has multiple plausible matches in $\mathbf{X_1}$, so that

Figure 1: Posterior means and credible intervals for accuracy metrics under the replication of simulation study from Sadinle 2017. For each level of overlap and each level of error, we have 100 paired sets of 500 records.

posterior match probability is split across more records. Since `BRL` removes previously matched records from the set of available records for a new match, it is able to retain higher confidence in matches. In practice, this partial estimate approach is not often used, so we remain confident in `fabl`'s performance. If one wishes to use partial estimates, `fabl` will possibly leave more linkages for the modeller to match by hand that would be left under `BRL`, but the decisions made by each method will have nearly equal accuracy.

8.3 Speed

To demonstrate speed, we generate comparison vectors from pre-specified distributions so that we can easily increase the size of the linkage problem. Distributions are meant to emulate the behavior of similarity scores across first name, last name, and day, month. For

Figure 2: Negative predictive value (NPV) and positive predictive value (PPV) on simulated datasets

example, $u^{\text{month}, 1} = P(\text{Records have same birth-month — Nonmatch}) = \frac{1}{12}$. For simplicity, we consider only exact matching, so a vector $(1, 0)$ corresponds to agreement and $(0, 1)$ to disagreement. We simulate these data for different values of n_A and n_B, and compare the run-time of `fabl` against `BRL`. Note that the number of unique patterns P is bounded above by $2^5 = 32$, a bound which is consistently attained in the larger simulations.

We see that at low data size, `BRL` outperforms, but that `fabl` is significantly faster at handling larger data. In particular, run-time for `BRL` seems to grow quadratically (or linearly with the size of both A and B) while run-time for `fabl` seems to grow linearly (in the size of B).

The above discussion suggests that for fixed n_B, computation time should remain mostly constant with growing n_A. Simulation study suggests that this is true. In the plot below, fixing $n_B = 500$, we see linear growth for the run-time under `BRL` as n_A increases, with much more static run-time under `fabl`. The slight increases in run-time that we do see are

	m	u
fname	$\left(\frac{19}{20}, \frac{1}{20}\right)$	$\left(\frac{1}{100}, \frac{99}{100}\right)$
lname	$\left(\frac{19}{20}, \frac{1}{20}\right)$	$\left(\frac{1}{100}, \frac{99}{100}\right)$
day	$\left(\frac{19}{20}, \frac{1}{20}\right)$	$\left(\frac{1}{30}, \frac{29}{30}\right)$
month	$\left(\frac{19}{20}, \frac{1}{20}\right)$	$\left(\frac{1}{12}, \frac{11}{12}\right)$
year	$\left(\frac{19}{20}, \frac{1}{20}\right)$	$\left(\frac{1}{12}, \frac{11}{12}\right)$

Table 2: Distributions used for m and u probabilties in simulation studies

Figure 3: Run-time for BRL and fabl to run 1000 Gibbs iterations, including hashing step for fabl, for increasing values of both n_A and n_B. We see near quadratic growth in runtime for BRL, and near linear growth for fabl.

due primarily to the hashing step, which again can be run in parallel for large data.

We note here that `BRL` is coded in C, which makes for unfair comparison against `fabl`, currently only built in R. Additionally, although `fabl` is amenable to parallelization, this simulation was run on a single core. Running `fabl` in C++ with paralellization for the hashing step and sampling the matching status of the record pairs should lead to even more computational gains.

8.4   Scaling to a Larger Simulation Study

In this section, we investigate scaling to larger data sets using our proposed methodology.

Figure 4: Run-time for BRL and fabl to run 1000 Gibbs iterations, including hashing step for fabl, with increasing n_A, and n_B fixed at 500. We see linear growth in runtime for BRL, and near constant runtime for fabl.

Specifically, we demonstrate how partitioning and hashing the data through `fabl` method allows us undertake significantly larger linkage tasks. We concatenate two sets of 40 simulated datasets from the original simulation studies of to create two larger files, each of 20,000 records (40,000 records total). Under standard Fellegi Sunter procedures, this would require 400,000,000 comparison vectors, each consisting for four integers, resulting a final comparison matrix about 6.4 GB in size. Under default settings on a personal computer, `R` does not allow storing an object this size, so `BRL` is not usable on this linkage task. Through `fabl` however, using just one machine, we sequentially compare records across chunks, hash results, and then synthesize summary statistics for the entire simulation.

For simplicity, we partition one dataset into 20 smaller chunks, and leave the second dataset fully intact. We compare records, hash results, and then synthesize summary statistics for all 20 chunk comparisons. The resulting data object is now only 90 MB, about

1% the size of the object required under the standard method. Executed sequentially, these comparisons and the Gibbs sampler take about one hour to run. However, this could be substantially sped up using distributed computing.

Turning to evaluation metrics, this simulation achieved 96.5% recall and 97.7% precision, with an overall F-measure of 97.1% F-measure. The reader will note that this slightly worse performance than witnessed in the smaller simulation studies; this is expected because it is naturally more difficult to link more files with the same amount of information. With more linkage fields `fabl` maintains the high accuracy seen above.

9 APPLICATION TO CIVILIAN CASUALTIES FROM THE EL SALVADORAN CIVIL WAR

In this section, we revisit an application to the El Salvadoran Civil War, where we estimate the number of documented identifiable deaths (DID). Though the data files used here are small, this study shows how the computational complexity of `fabl` depends on the number of unique agreement patterns, and how significant computational gains can be achieved by simplifying the construction of the comparison vectors. This case study also considers the ramifications of independently sampling Z_j rather than strictly enforcing one-to-one matching as done by `BRL`.

9.1 El Salvadoran Civil War

For Beka: You wrote that you wanted more detail here, and I added some of it (with some citations). I believe I have all the important informatoin, just more concise. If I included Sadinle's level of detail, I worried about plagarism.

The country of El Salvador was immersed in civil war from 1980 to 1991, and throughout the time, several organizations attempted to document casualties of the conflict. When estimating the total number of casualties, one cannot simply sum the numbers recorded by each organization, as it is likely that the same individuals are recorded in multiple casualty lists. This is an important step in estimating the total number of deaths from a conflict

using multiple systems estimation (Lum et al., 2013) Thanks to the Human Rights Data Analysis Group (HRDAG), we utilize two different sources, namely, El Rescate - Tutela Regal (ERTL) and the Salvadoran Human Rights Commission (CDHES, by its acronym in Spanish). The ERTL dataset consists of digitized denunciations that had been published throughout the conflict, and additional investigation was required before the ERTL recorded them as human rights abuses, giving us confidence in its accuracy Howland (2008). The CDHES dataset consists of casualties that had been reported directly to the organization, and later digitized; since reports were taken soon after the events occured, this dataset is also thought to be fairly reliable. Ball (2000).

There are several challenges with working with such data. Firstly, both datasets have been automatically digitized, which inherently leads to some degree of typographical error. Secondly, the CDHES records are all second hand accounts reported by individuals, which can result in additional errors. Lastly, the only fields recorded are given name, last name, date of death, and place of death; it is relatively common for a parent and child to share the same given name, resulting in indistinguishable records for two different individuals. This last point nearly breaks the earlier mentioned assumption that there are no duplicates within files, and reveals a key difference between `BRL` and our proposed method.

In this paper, we utilize records that have non-missing entries for given and last name, which results in $n_1 = 4420$ files in CHDES and $n_2 = 1323$ files in ERTL. We standardized names to account for common misspellings in the Spanish language. The string distance used is a modified Levenstein distance to account for the fact that second names are often omitted. Place of death is recorded by municipality and department within that municipality; however, since department was missing in 95% of records in CHDES and 80% of records in ERTL, this was excluded from our analyses. Thus, we conduct record linkage using given name, last name, municipality, and day, month, and year of death. We again use flat priors for the **m** and **u** parameters. To our knowledge, we have followed the same guidelines/settings in Sadinle (2017).

To mirror the original implementation of Sadinle (2017), we construct the comparison vectors using four levels of agreement for each field, according to the thresholds provided in

Fields	Similarity	Level of Disagreement			
		1	2	3	4
First and Last Name	Modified Levenstein	0	(0, .25]	(.25, .5]	(5, 1]
Year of Death	Absolute Difference	0	1	2	3+
Month of Death	Absolute Difference	0	1	2-3	4+
Day of Death	Absolute Difference	0	1-2	3-7	8+
Municipality	Binary	Agree	Disagree		

Table 3: Construction of comparison vectors for El Salvador data resembling original implementation from (Sadinle, 2018). This set up leads to 2048 possible agreement patterns in total

Fields	Similarity	Level of Disagreement		
		1	2	3
First and Last Name	Modified Levenstein	0	(0, .25]	(.25, 1]
Year of Death	Binary	Agree	Disagree	
Month of Death	Binary	Agree	Disagree	
Day of Death	Absolute Difference	0	1	2+
Municipality	Binary	Agree	Disagree	

Table 4: Construction of comparison vectors for El Salvador for increased speed under `fabl`. This set up leads to 216 possible agreement patterns in total

Table 3. Gibbs sampling for these comparison vectors took 422 seconds for our proposed method, `fabl`, and 240 for `BRL`. However, we observed that posterior distributions of several levels of the **m** and **u** parameters were nearly identical, and that of the $4^5 \times 2 = 2048$ possible agreement patterns, only 1173 are realized in the data. This leads us to believe that such high number of agreement levels creates unnecessary distinctions in the data and makes the comparison vectors less interpretable.

Therefore we re-ran our analysis with fewer agreement levels for each field according to Table 4, and obtained analogous results. With 216 possible agreement patterns, 159 were realized in the data, and our proposed method, `fabl` became much faster, finishing in 61 seconds. Meanwhile `BRL` took 239 seconds, relatively unchanged from the first implementation. This demonstrates the way that the computational complexity of our method depends on the number of unique agreement patterns, and how significant computational gains can be made by simplifying the construction of the comparison vectors. The estimates of the **m** parameters under each method are very similar, as shown in Figure 6.

Through `fabl`, we arrive at a Bayes estimate of 179 individuals recorded in both datasets. We calcuate posterior samples of size of the overlap across files by summing the number of matches found in each iteration of the Gibbs sampler, and subtracting the number of matches that violate one-to-one matching. The posterior 95% credible interval for the overlap across files is (206, 238), indicating that the Bayes estimate is quite conservative. This is because there a large number of records in ERTL for which there are multiple plausible matches in CDHES; `fabl` recognizes that a match exists among the several options, but is unable to definitely declare a specific pair a match in the Bayes estimate. We also compute a partial estimate of the linkage structure, using $\theta_{10} = \theta_{01} = 1$, $\theta_{11} = 2$, and $\theta_R = 0.1$ as in the simulation study. Here, the Bayes estimate provides 136 matches of which the model is quite confident, and 175 records to verify manually; this means that after clerical review, the number of individuals replicated across datasets would fall in the interval (136, 311), encapsulating the posterior credible interval. More or fewer records could be indentified for clerical review by decreasing or increasing θ_R.

We see similar results under `BRL`, with a Bayes estimate of 181 individuals recorded in

both datasets, a posterior 95% credible interval of (211, 244), and a range of (140, 294) after the partial estimate and clerical review.


Figure 5: Posterior distribution of overlap across the two files. The solid lines show the Bayes estimate for the amount of overlap, and the dashed line is the Bayes estimate under fabl before resolving violations of one-to-one matching

Violations of One-to-One Matching   As noted previously, the application exhibits many situations in which bipartite matching is difficult. To explore further, we calculate the number of matchings throughout the `fabl` Gibbs sampler that violate our one-to-one assumption, and find on average 26.17 matchings in violation. These matchings are impossible under the `BRL` Gibbs sampler, so examining these more closely can provide insight how each model handles situations in which one file in X_1 has multiple plausible matches in X_2.

Figure 7 illustrates an example. Note that these records present a near violation of our assumption that there are no duplications within files; we continue to assume that

Figure 6: Posterior estimates of m parameters with 95% credible intervals

records 825 and 826 in ERTL correspond to different individuals, but their records are nearly identical. In addition, using the modified Levenstein distance, the comparison vectors $\gamma_{2776,825}$ and $\gamma_{2776,826}$ are exactly identical.

Figure 7: Example of linkage situation with multiple plausible matches

Figure 8 shows the values that Z_{825} and Z_{826} take on throughout the Gibbs sampler, and demonstrates how each method handles this situation. Under `fabl`, both records in $\boldsymbol{X}_2$ match to the same record in $\boldsymbol{X}_1$ throughout the Gibbs process, creating consistent violations of one-to-one matching. Under `BRL`, the Gibbs process creates one matching configuration and stays there for a while. However, if one pair "unmatches," then the other record has a chance to latch on. Then, the Gibbs process is stuck with that matching status for a while, resulting in a Gibbs process with poor mixing. In addition, `fabl` allows the modeler to inspect records with multiple plausible matches, and if they desire, to then choose the record pairing with the highest posterior probability. `BRL` in contrast, in strictly

44

enforcing one-to-one matching throughout the sampler, can lead to situations where none of the plausible matches reach the threshold to be identified through the Bayes estimate. In short, we have pointed out why one may not wish to imbed this constraint in a model, and proposed a simple alternative that performs about the same as BRL, while scaling to much larger data sets.
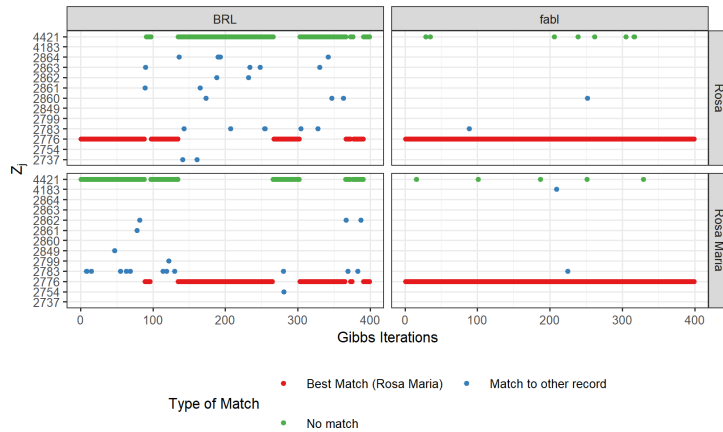

Figure 8: Gibbs sampling in situation with multiple plausible matches.

10   DISCUSSION

We have presented a method for Bayesian record linkage that is feasible for large datasets. In particular, our hashing procedure and model assumptions allow for a linkage procedure whose computational complexity does not scale with the size of the larger dataset, making this method particularly powerful in linking records when one datafile is substantially smaller than the other.

In our case study, we included an exploration of how to conduct record linkage when modeling assumptions are not met in practice. We explored "one-to-many" scenarios in which one record in A has multiple plausible matches in B, and showed how both fabl and BRL demonstrated undesirable qualities. Other issues arise under "many-to-one" scenarios, where one record in B has multiple plausible matches in A, and "many-to-many" scenarios in which there is duplication both across and within datasets. Tuning fabl for use in these

45

scenarios is one potential avenue for future work.

11 APPENDIX

Figure 9: Traceplot for number of matches found across datasets in El Salvador case study

REFERENCES

Ball, P. (2000), "The Salvadoran Human Rights Commission: Data Processing, Data Representation, and Generating Analytical Reports," in *Making the Case: Investigating Large Scale Human Rights Violations Using Information Systems and Data Analysis*, eds. P. Ball, H. F. Spirer, and L. Spirer, pp. 15–24, American Association for the Advancement of Science.

Belin, T. R. and Rubin, D. B. (1995), "A Method for Calibrating False-Match Rates in Record Linkage," *Journal of the American Statistical Association*, 90, 694–707.

Bilenko, M. and Mooney, R. (2006), "Riddle: Repository of Information on Duplicate Detection, Record Linkage, and Identity Uncertainty," Online; retrieved July 29, 2020.

Figure 10: Traceplot for m parameter in El Salvador case study

Bilenko, M., Kamath, B., and Mooney, R. J. (2006), "Adaptive blocking: Learning to scale up record linkage," *Proceedings - IEEE International Conference on Data Mining, ICDM*, pp. 87–96.

Christen, P. (2008), "Automatic Record Linkage Using Seeded Nearest Neighbour and Support Vector Machine Classification," in *Proceedings of the 14th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, p. 151–159, New York, NY, USA, Association for Computing Machinery.

Christen, P. (2012), "A survey of indexing techniques for scalable record linkage and deduplication," *IEEE Transactions on Knowledge and Data Engineering*, 24, 1537–1555.

Dalzell, N. M. and Reiter, J. P. (2018), "Regression modeling and file matching using possibly erroneous matching variables," *Journal of Computational and Graphical Statistics*, 27, 728–738.

Elmagarmid, A. K., Ipeirotis, P. G., and Verykios, V. S. (2007), "Duplicate Record Detection: A Survey," *IEEE Transactions on Knowledge and Data Engineering*, 19, 1–16.

Enamorado, T., Fifield, B., and Imai, K. (2019a), "Using a Probabilistic Model to Assist Merging of Large-Scale Administrative Records," *American Political Science Review*, 113, 353–371.

Figure 11: Traceplot for u parameter in El Salvador case study

Enamorado, T., Fifield, B., and Imai, K. (2019b), "Using a Probabilistic Model to Assist Merging of Large-Scale Administrative Records," *American Political Science Review*, 113, 353–371.

Fair, M. (2004), "Generalized record linkage system–Statistics Canada's record linkage software," *Austrian Journal of Statistics*, 33, 37–53.

Fellegi, I. P. and Sunter, A. B. (1969), "A Theory for Record Linkage," *Journal of the American Statistical Association*, 64, 1183–1210.

Fortini, M., Liseo, B., Nuccitelli, A., and Scanu, M. (2001), "On Bayesian Record Linkage," *Research in Official Statistics*, 4, 185–198.

Fortunato, S. (2010), "Community detection in graphs," *Physics Reports*, 486, 75–174.

Gill, L. and Goldacre, M. (2003), "English national record linkage of hospital episode statistics and death registration records," *Report to the department of health*.

Gutman, R., Afendulis, C. C., and Zaslavsky, A. M. (2013), "A Bayesian Procedure for File Linking to Analyze End-of-Life Medical Costs," *Journal of the American Statistical Association*, 108, 34–47.

Herzog, T. N., Scheuren, F. J., and Winkler, W. E. (2007), *Data quality and record linkage techniques*, Springer Science & Business Media.

Howland, T. (2008), "How El Rescate, a Small Nongovernmental Organization, Contributed to the Transformation of the Human Rights Situation in El Salvador," *Human Rights Quarterly*, 30, 703–757.

Jaro, M. A. (1989), "Advances in Record-Linkage Methodology as Applied to Matching the 1985 Census of Tampa, Florida," *Journal of the American Statistical Association*, 84, 414–420.

Larsen, M. (2005a), "Hierarchical Bayesian Record Linkage Theory," .

Larsen, M. D. (2005b), "Advances in Record Linkage Theory: Hierarchical Bayesian Record Linkage Theory," in *Proceedings of the Joint Statistical Meetings, Section on Survey Research Methods*, pp. 3277–3284, The American Statistical Association.

Larsen, M. D. (2012), "An Experiment with Hierarchical Bayesian Record Linkage," arXiv e-prints, arxiv:1212.5203.

Larsen, M. D. and Rubin, D. B. (2001), "Iterative Automated Record Linkage Using Mixture Models," *Journal of the American Statistical Association*, 96, 32–41.

Liseo, B. and Tancredi, A. (2011), "Bayesian estimation of population size via linkage of multivariate Normal data sets," *Journal of Official Statistics*, 27, 491–505.

Lum, K., Price, M. E., and Banks, D. (2013), "Applications of multiple systems estimation in human rights research," *The American Statistician*, 67, 191–200.

Marchant, N. G., Steorts, R. C., Kaplan, A., Rubinstein, B. I. P., and Elazar, D. N. (2019), "d-blink: Distributed End-to-End Bayesian Entity Resolution," *arXiv e-prints*, arxiv:1909.06039.

Matsakis, N. E. (2010), "Active Duplicate Detection with Bayesian Nonparametric Models," Ph.D. thesis, Massachusetts Institute of Technology.

Newcombe, H. B. and Rhynas, P. O. W. (1962), "Child Spacing Following Stillbirth and Infant Death," *Eugenics Quarterly*, 9, 25–35.

Newcombe, H. B., Kennedy, J. M., Axford, S. J., and James, A. P. (1959), "Automatic Linkage of Vital Records," *Science*, 130, 954–959.

Sadinle, M. (2017), "Bayesian Estimation of Bipartite Matchings for Record Linkage," *Journal of the American Statistical Association*, 112, 600–612.

Sadinle, M. (2018), "Bayesian Propagation of Record Linkage Uncertainty Into Population Size Estimation of Human Rights Violations," *The Annals of Applied Statistics*, 12, 1013–1038.

Shan, M., Thomas, K., and Gutman, R. (2020), "A Bayesian Multi-Layered Record Linkage Procedure to Analyze Functional Status of Medicare Patients with Traumatic Brain Injury," *arXiv e-prints*, arxiv:2005.08549.

Steorts, R. (2013), "Divergence losses under benchmarking for small area estimation," *Working Paper*.

Steorts, R. C. (2015), "Entity Resolution with Empirically Motivated Priors," *Bayesian Analysis*, 10, 849–875.

Steorts, R. C., Hall, R., and Fienberg, S. E. (2016), "A Bayesian Approach to Graphical Record Linkage and Deduplication," *Journal of the American Statistical Association*, 111, 1660–1672.

Steorts, R. C., Tancredi, A., and Liseo, B. (2018), "Generalized Bayesian Record Linkage and Regression with Exact Error Propagation," in *International Conference on Privacy in Statistical Databases*, pp. 297–313, Springer.

Tancredi, A. and Liseo, B. (2011), "A Hierarchical Bayesian Approach to Record Linkage and Size Population Problems," *Annals of Applied Statistics*, 5, 1553–1585.

Tancredi, A. and Liseo, B. (2015), "Regression analysis with linked data: problems and possible solutions," *Statistica*, 75, 19–35.

Tancredi, A., Steorts, R., and Liseo, B. (2020), "A Unified Framework for De-Duplication and Population Size Estimation (with Discussion)," *Bayesian Analysis*, 15, 633–682.

Ventura, S. L., Nugent, R., and Fuchs, E. R. (2015), "Seeing the Non-Stars:(Some) Sources of Bias in Past Disambiguation Approaches and a New Public Tool Leveraging Labeled Records," *Research Policy*, 44, 1672–1701.

Wagner, D., Lane, M., et al. (2014), "The person identification validation system (PVS): applying the Center for Administrative Records Research and Applications'(CARRA) record linkage software," Tech. rep., Center for Economic Studies, US Census Bureau.

Winkler, W. E. (1988), "Using the EM Algorithm for Weight Computation in the Fellegi-Sunter Model of Record Linkage," in *Proceedings of the Section on Survey Research Methods*, pp. 667–671, American Statistical Association.

Winkler, W. E. and Thibaudeau, Y. (1991), *An Application of the Fellegi-Sunter Model of Record Linkage to the 1990 US Decennial Census*.

Zanella, G., Betancourt, B., Wallach, H., Miller, J., Zaidi, A., and Steorts, R. C. (2016), "Flexible Models for Microclustering with Application to Entity Resolution," in *Proceedings of the 30th International Conference on Neural Information Processing Systems*, NIPS'16, pp. 1425–1433, NY, USA, Curran Associates Inc.