

Unveiling parl

Brian Kunder

Overview

- ▶ Computational Speed-ups
- ▶ Next Step

- ▶ I have been playing around with calling my method “Parallelized Record Linkage in R”, or parlr for short.
- ▶ Broadly speaking, I use two strategies for computational speedup:
 - ▶ Performing operations on the **unique agreement patterns** rather than the individual record pairs. Rigorously, sums over the unique patterns are *sufficient statistics*.
 - ▶ Relaxation of one to one matching to allow for parallel sampling of Z

Together, these have made **huge** improvements on code

Identifying Unique Patterns

- ▶ Enamorado uses the following hashing function to identify unique agreement patterns. I've also figured out how to do this efficiently in `dplyr`.

$$H(i, j) = \sum_{f=1}^F \mathbf{1}_{\gamma_f(i, j) > 0} 2^{\gamma_f(i, j) + \mathbf{1}_{f > 1} \times \sum_{a=1}^{f-1} (L_a - 1)}$$

- ▶ I identify the unique patterns, the number of each of those patterns in the data as a whole, and the number of those patterns for each record $j \in B$.
- ▶ Define P to be the number of unique agreement patterns, and note it is bounded by $\prod_{f=1}^F L_f$. This *does not* grow with the size of the data.

Calculating FS weights

- ▶ Within a Gibbs iteration, all records with the same agreement pattern will have the same FS weight. So instead of calculating the weight for each of $n_A \times n_B$ records, I only need to do it for P unique patterns.
- ▶ This step is taken directly from `fastlink`.

Identifying Unique Patterns

- ▶ `fastlink` doesn't have posterior updates for Dirichlet distributions, so I came up with this step on my own.
- ▶ Each agreement pattern corresponds to a particular type of contribution to the posteriors of the m and u parameters.
- ▶ To update m , I used to sum over each variable for all N record pairs. Now, I simply identify the number of matchings in the Z vector corresponding to each of P patterns.
- ▶ The nonmatches are even easier: just subtract this number from the total present in the data!

Breaking up the Z sampler

- ▶ Four observations
 - ▶ Sampling many objects with varying probabilities difficult
 - ▶ Sampling *fewer* objects with varying probabilities is easy
 - ▶ Sampling many objects with *uniform* probabilities is also easy
 - ▶ For a record $j \in B$, each record $i \in A$ with a given agreement pattern has the same probability of matching. That is, the unique patterns form *equivalence classes*, and within a class, each record is equally likely.

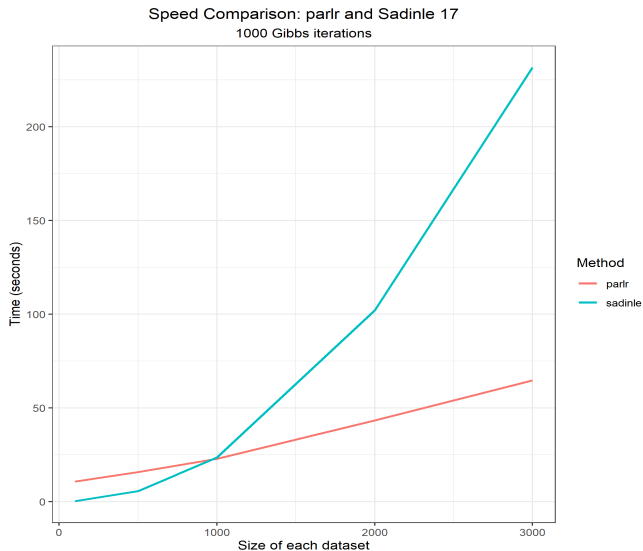
Breaking up the Z sampler

- ▶ Rather than sampling Z_j sequentially from N many records, I can sample it in parallel from P many equivalence classes of unique patterns, each one multiplied by its prevalence for record j .
- ▶ After sampling the unique pattern, I sample the record associated with that pattern uniformly at random.
- ▶ Since all calculations within the Gibbs Sampler depend only on the unique pattern, and since my model samples the Z_j independently, I can sample for the records *at the end of the entire Gibbs procedure*

Summary

- ▶ We have *almost entirely* removed dependence on N from our Gibbs procedure, and replaced them with steps that depend on P .
- ▶ Only major hurdle with N is calculating the comparison vectors, which is just unavoidable in the Fellegi Sunter Framework. However, whatever Ted does for `fastlink` would work here too.
- ▶ Filling in the Gibbs sampler post-hoc technically depends on N , but its a very inexpensive step that only has to be done once

Speed up



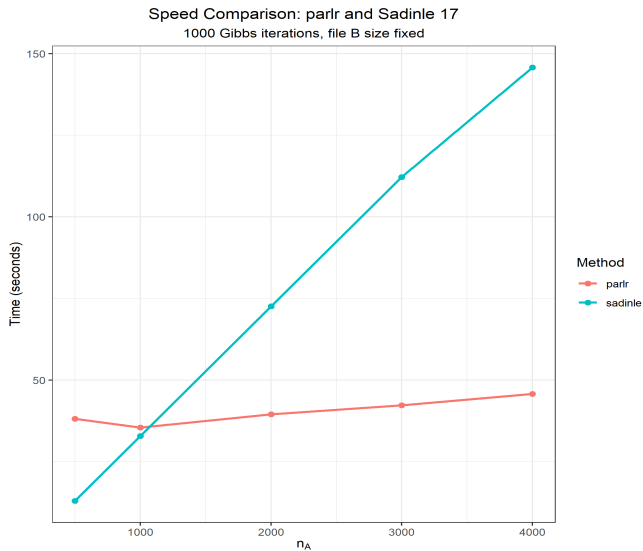
Speed up

- ▶ When n_A and n_B are small, the computational savings are minimal, and Sadinle's BRL runs faster just because its in C.
- ▶ As the datasets grow bigger, computation for BRL grows *quadratically*, while `parlr` seems to grow *linearly*
- ▶ The plot shown is with **no parallel computing**
- ▶ If my method was programmed in C, I would expect even greater computational savings


More on Speed Up

- ▶ Within the Gibbs sampler, calculations only depend on n_B (the number samples I need to take to fill a Z vector), and P (the number of unique patterns). n_A only influences the scale of the counts for the unique patterns.
- ▶ n_A only comes to play in forming the comparison vectors and in back-filling the samples of Z at the end of the sampler.
- ▶ This suggests that the sampler should be *almost constant* in n_A (ie, the Gibbs sampling for a 4000×500 problem should not be substantially longer than for a 500×500 problem)
- ▶ Particularly useful for linking a smaller dataset to a much larger dataset. One example would be the DNC data in BRACS!

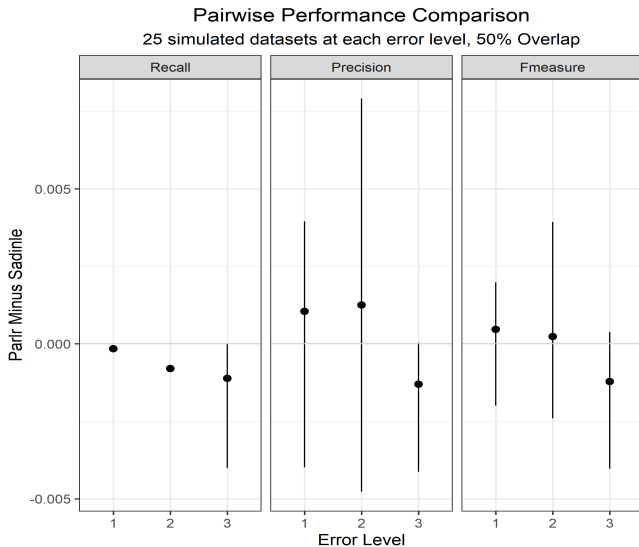
More on Speed Up



Accuracy



Accuracy



Accuracy

- ▶ Overall, the `parlr` method performs comparably to Sardinia17. There is slight evidence that recall is systematically lower for `parlr` at high error levels, but it is very slight. I'll need to run this at different level of overlap to explore more.

Limitations

- ▶ The original FS method is $O(n_A \times n_B)$ hard. With hashing methods, `fastlink` makes the problem only $O(P)$ hard, since record pairs are considered independently and classification is based *only* on the unique pattern.
- ▶ Hashing methods in my model make the problem $O(n_B \times P)$ hard. No matter what, I need to sample n_B many times to fill a linkage vector Z . In this sense, `fastlink` is inherently faster than this method
- ▶ However, method takes into consider much more of the dependency structure in the data, making this computational burden worthwhile
- ▶ **Note:** If this line of thinking is useful for the paper, I'll need some help making this rigorous!

Next Steps

- ▶ I have not yet implemented any parallelization. We will cover this soon in Shawns class. The steps above however provide *the vast majority* of the computational speed up.
- ▶ Later on, I'll be able to use this same hashing function for linkage clusters by replacing 2 with the k^{th} prime number for the k^{th} linkage cluster. (For Jody's paper, I'll present just the standard method)

Next Steps

- ▶ This model and these computational savings really inform everything everything I'm doing. Apart from the BRACS paper with Jody, I'm really going to focus on getting this done
- ▶ Larger versions of the simulations above, with different levels of overlap.
- ▶ **Question:** What other kinds of simulations or evaluations would I need?
- ▶ **Question:** Is this still a good fit for Bayesian Analysis? Or should it go to a journal more concerned with computation and scalability?